

# Paper reproduction project Segnet - Group 26

Hao Li 5008891, Zixuan Wan 5059747, Nouredine Begga 4384261

April 20, 2020

## 0.1 Our take on the Impact of Illumination-Invariant Image Pre-transformation for Contemporary Automotive Semantic Scene Understanding.

In 2018, Naif Alshammari, Samet Akcay and Toby P. Breckon published their paper On the Impact of Illumination-Invariant Image Pre-transformation for Contemporary Automotive Semantic Scene Understanding. The authors of this paper realized that illumination changes in outdoor environments under non-ideal weather conditions have a negative impact on automotive scene understanding and segmentation performance. The paper presents an evaluation of illumination-invariant image transforms on the CamVid dataset, to see if this improves the state of the art in scene understanding performance.

As part of TU Delft's CS4240 Deep Learning course, we — Nouredine Begga, Hao Li and Zixuan Wan— attempt to give the reproduction of the results achieved in said paper a try. That is, we try to develop an implementation of the Deep Fully Convolutional Neural Network SegNet and preprocess the CamVid dataset with the illumination invariant image representation. More specifically, our goal is to achieve similar results as Table 1 of the paper (which is copied below).

Method	Sky	Building	Pole	Road	Pavement	Tree	SignSymbol	Fence	Car	Pedestrian	Bicyclist	Class avg.	Global acc.	mIoU	Precision	Recall
Original (RGB) [1]	0.73	0.846	0.33	0.87	<b>0.91</b>	0.76	0.43	0.41	0.73	0.60	0.11	0.61	0.807	0.46	<b>0.70</b>	0.61
$\mathcal{I}_{\text{Alvarez}}$ [7]	0.67	0.73	0.22	0.69	0.67	0.75	0.35	0.28	0.63	0.26	0.017	0.46	0.68	0.33	0.46	0.48
$\mathcal{I}_{\text{Maddern}}$ [6]	0.92	0.80	0.20	0.932	0.53	0.62	0.38	0.17	0.61	0.51	0.07	0.54	0.78	0.40	0.64	0.65

TABLE 1: Quantitative results are shown as accuracy of the CNN SegNet approach on CamVid test data for RGB and two methods.

The Cambridge-driving Labeled Video Database (CamVid) is the first collection of videos with object class semantic labels, complete with metadata. The database provides ground truth labels that associate each pixel with one of 32 semantic classes.



Figure 1: Obtained RGB images and their respective ground truth class labels

## 0.2 What is a illumination invariant image representation?

An illumination invariant image representation is a colour representation computed from RGB that removes (or minimises) scene colour variations due to varying scene lighting conditions. This technique was introduced as an intrinsic image to represent the illumination invariant and intrinsic properties in the image [1] with illumination transforms generally computed with reference to the physical properties behind the capture and the presence of colour within the space. In most literature where a type of illumination invariance is applied had as an objective to remove shadows, and to improve scene classification and

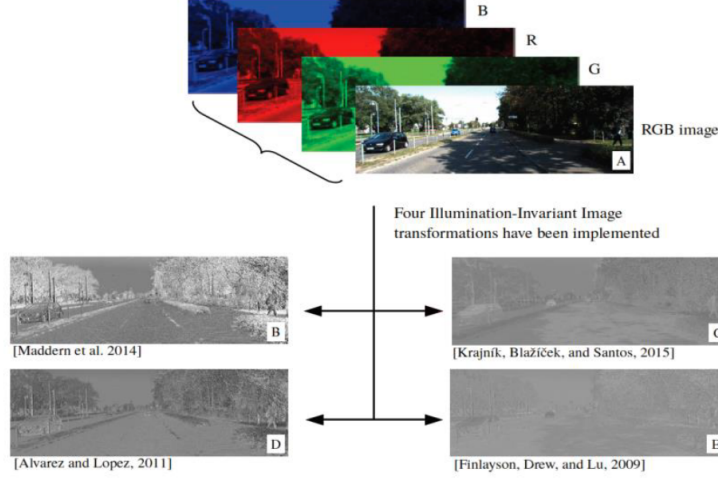


Figure 2: An example of an RGB image followed by four different illumination invariant images, where all the illumination variations such as shadows are significantly reduced within the scenes.

segmentation.

For

our reproduction we will implement two illumination invariance methods and compare these with the results of the authors. For our reproduction we will use the methods of Maddern et al.[2] and Alvarez et al.[3].

## 0.3 illumination invariance method - Alvarez et al.

An illumination-invariant image  $I$  is a single channel image calculated by combining the three RGB colour channels in the image  $IRGB \in \{IR, IG, IB\}$ . To compute the illumination invariant images, we use a single channel feature space  $I$  combined with three linear sensors  $\{R, G, B\}$  as follows for the method of Alvarez et al:

$$I_{\text{Alvarez}} = \cos(\theta) \log_{\text{approx}} \left( \frac{I_R}{I_B} \right) + \sin(\theta) \log_{\text{approx}} \left( \frac{I_G}{I_B} \right)$$

Where  $I_R, [G, IB$  are the tree RGB channels,  $\theta \in \{0...180\}$ , and  $\log_{\text{approx}} 0$  is the logarithmic approximation for  $t \in \{R, G, B\}$ , which is computes as follows:


$$\log_{\text{approx}(x)} = \alpha \left( \left( x^{\frac{1}{\alpha}} \right) - 1 \right)$$

Where  $x$  is the value from dividing the two channels and  $\alpha = 5000$ . After evaluating images we decide projection angle  $\theta = 135$  degree.## illumination invariance method - Maddern et al. To compute the image for this method, we again convert the 3-channel floating point RGB image into corresponding illumination invariant image as follows:

$$I_{\text{Maddern}} = 0.5 + \log(I_G) - \alpha \log(I_B) - (1 - \alpha) \log(I_R)$$

Where  $\alpha = 0.48$ . This illumination-invariant approach was proposed to improve visual localization, mapping and scene classification for autonomous road vehicles.

Our results for image processing:



#### 0.4 How will we evaluate the performance of using these illumination invariant image representations?

We will evaluate the performance of automotive scene understanding and segmentation using the SegNet [4] CNN architecture (Figure 3) with the two aforementioned illumination-invariant transformations.

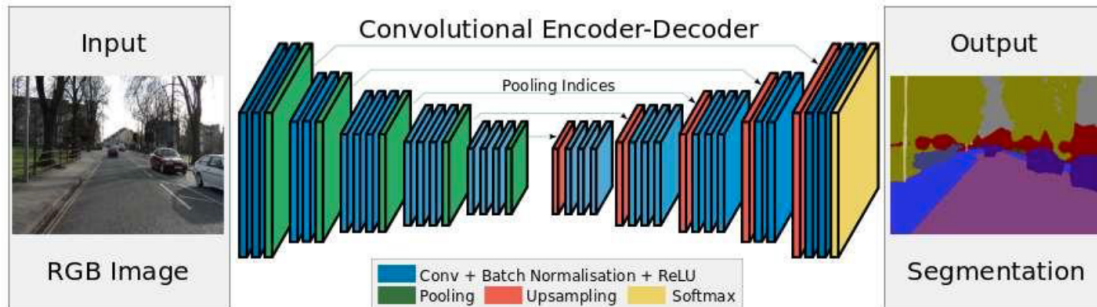


Figure 3: Architecture of the SegNet Convolutional Neural Network

As mentioned before we will use the CamVid dataset with different pixel classes for the SegNet classification task (Figure 1 shows some examples). The authors of the paper used eleven classes: {sky, building, pole, road, pavement, tree, sign, fence, car, pedestrian, bicycle}. The dataset consists of 600 images in total, which we have divided into training, test and validation sets.

#### 0.5 Training the SegNet CNN model

The authors used a VGG16 [5] network pre-trained on the ImageNet [6] dataset, which is the encoder network within SegNet. An encoder network consists of convolution and pooling layers followed by a decoder network containing convolutional and upsampling layers. The authors have used Stochastic Gradient Descent (SGD) optimization, we however have chosen to use Adam optimization in the hopes of getting better results than the authors. Adam can be looked at as a combination of RMSprop and Stochastic Gradient Descent with momentum. It uses the squared gradients to scale the learning rate like RMSprop and it takes advantage of momentum by using moving average of the gradient instead of gradient itself like SGD with momentum[7]. We used the following hyperparameters: initial learning rate  $1 \times 10^{-3}$ , no weight decay  $5 \times 10^{-4}$  and momen-

tum 0.9. We train the model using the freely provided GPU from Google Colab. For training the neural network on the whole dataset, because of the computation capacity, we downsized original images with scale ratio 0.1. The image resolution then become 72x96, which is 1/5 smaller than the original data. The amount of images for train, test and validation sets are 447,58,196 respectively, which has approximate proportion of 0.6,0.1,0.3. Self-developed dataloader is used to separate, label processing and transform raw images.

## 0.6 Code implementation

```
[1]: # Import Libraries
from __future__ import print_function
import shutil
import os
import torch
import torchvision.transforms as transforms
import torchvision.datasets as datasets
import torchvision.models as models
import torch.nn as nn
import torch.optim as optim
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
%matplotlib inline
import math
from collections import OrderedDict
import torch.nn.functional as F
import pprint
import torch.utils.data as data
import sys
sys.path.append('./')
import utils
import random
```

```
[2]: # Define dataloader class
class CamVid(data.Dataset):
    """CamVid dataset loader where the dataset is arranged as in
    https://github.com/alexgkendall/SegNet-Tutorial/tree/master/CamVid.
    Keyword arguments:
    - root_dir (`string`): Root directory path.
    - mode (`string`): The type of dataset: 'train' for training set, 'val'
    for validation set, and 'test' for test set.
    - transform (`callable`, optional): A function/transform that takes in
    an PIL image and returns a transformed version. Default: None.
    - label_transform (`callable`, optional): A function/transform that takes
    in the target and transforms it. Default: None.
    - loader (`callable`, optional): A function to load an image given its
    path. By default ``default_loader`` is used.
```

```

"""
# Training dataset root folders
train_folder = 'train'
train_lbl_folder = 'trainannot'

# Validation dataset root folders
val_folder = 'val'
val_lbl_folder = 'valannot'

# Test dataset root folders
test_folder = 'test'
test_lbl_folder = 'testannot'

# Images extension
img_extension = '.png'

def __init__(self,
              root_dir,
              mode='train',
              transform=None,
              label_transform=None,
              loader=utils.pil_loader):

    self.root_dir = root_dir
    self.mode = mode
    self.transform = transform
    self.label_transform = label_transform
    self.loader = loader

    if self.mode.lower() == 'train':
        # Get the training data and labels filepaths
        self.train_data = utils.get_files(
            os.path.join(root_dir, self.train_folder),
            extension_filter=self.img_extension)

        self.train_labels = utils.get_files(
            os.path.join(root_dir, self.train_lbl_folder),
            extension_filter=self.img_extension)
    elif self.mode.lower() == 'val':
        # Get the validation data and labels filepaths
        self.val_data = utils.get_files(
            os.path.join(root_dir, self.val_folder),
            extension_filter=self.img_extension)

        self.val_labels = utils.get_files(
            os.path.join(root_dir, self.val_lbl_folder),
            extension_filter=self.img_extension)

```

```

elif self.mode.lower() == 'test':
    # Get the test data and labels filepaths
    self.test_data = utils.get_files(
        os.path.join(root_dir, self.test_folder),
        extension_filter=self.img_extension)

    self.test_labels = utils.get_files(
        os.path.join(root_dir, self.test_lbl_folder),
        extension_filter=self.img_extension)
else:
    raise RuntimeError("Unexpected dataset mode. "
                       "Supported modes are: train, val and test")

def __getitem__(self, index):
    """
    Args:
        - index (`int`): index of the item in the dataset
    Returns:
        A tuple of ``PIL.Image`` (image, label) where label is the ground-truth
        of the image.
    """
    if self.mode.lower() == 'train':
        data_path, label_path = self.train_data[index], self.train_labels[
            index]
    elif self.mode.lower() == 'val':
        data_path, label_path = self.val_data[index], self.val_labels[
            index]
    elif self.mode.lower() == 'test':
        data_path, label_path = self.test_data[index], self.test_labels[
            index]
    else:
        raise RuntimeError("Unexpected dataset mode. "
                           "Supported modes are: train, val and test")

    img, label = self.loader(data_path, label_path)
    if self.transform is not None:
        img = self.transform(img)

    if self.label_transform is not None:
        label = self.label_transform(label)
    label = label.permute(1,2,0)*255
    label_np = label.numpy()
    label = rgb_to_label(label_np, colormap=color_encoding)
    label = torch.LongTensor(label)
    return img, label

def __len__(self):

```

```

"""Returns the length of the dataset."""
if self.mode.lower() == 'train':
    return len(self.train_data)
elif self.mode.lower() == 'val':
    return len(self.val_data)
elif self.mode.lower() == 'test':
    return len(self.test_data)
else:
    raise RuntimeError("Unexpected dataset mode. "
                       "Supported modes are: train, val and test")

```

```

[3]: # Default encoding for pixel value, class name, and class color
class_name = {0: 'sky',
              1: 'building',
              2: 'pole',
              3: 'road_marking',
              4: 'road',
              5: 'pavement',
              6: 'tree',
              7: 'sign_symbol',
              8: 'fence',
              9: 'car',
              10: 'pedestrian',
              11: 'bicyclist',
              12: 'unlabeled'}

color_encoding = {0: (128, 128, 128),
                  1: (128, 0, 0),
                  2: (192, 192, 128),
                  3: (255, 69, 0),
                  4: (128, 64, 128),
                  5: (60, 40, 222),
                  6: (128, 128, 0),
                  7: (192, 128, 128),
                  8: (64, 64, 128),
                  9: (64, 0, 128),
                  10: (64, 64, 0),
                  11: (0, 128, 192),
                  12: (0, 0, 0)}

def rgb_to_label(rgb_image, colormap = color_encoding):
    '''Function to one hot encode RGB mask labels
    Inputs:
        rgb_image - image matrix (eg. 256 x 256 x 3 dimension numpy ndarray)
        colormap - dictionary of color to label id

```



```

        Output: One hot encoded image of dimensions (height x width x
        →num_classes) where num_classes = len(colormap)
        '''
        num_classes = len(colormap)
        shape = rgb_image.shape[:2]
        encoded_image = np.zeros(shape,dtype=np.int8)
        for i, cls in enumerate(colormap):
            for x in range(encoded_image.shape[0]):
                for y in range (encoded_image.shape[1]):
                    if(np.all(rgb_image[x][y] == colormap[i])):
                        encoded_image[x][y] = i
        return encoded_image

def label_to_rgb(label, colormap = color_encoding):
    '''Function to decode encoded mask labels
    Inputs:
        onehot - one hot encoded image matrix (height x width x num_classes)
        colormap - dictionary of color to label id
    Output: Decoded RGB image (height x width x 3)
    '''
    output = np.zeros(label.shape[:2]+(3,))
    for k in colormap.keys():
        output[label==k] = colormap[k]
    return np.uint8(output)

```

```

[4]: # Specify transforms using torchvision.transforms as transforms

res_ratio = 0.1 # set image size

transformations = transforms.Compose([
    transforms.
    →Resize((int(res_ratio*720),int(res_ratio*960)),interpolation=Image.NEAREST),
    →#set resolution with nearest neighbor interpolation
    transforms.ToTensor() # Normalize data to be of values [0-1]
])

train = CamVid('./','train',transformations, transformations)
val = CamVid('./','val',transformations, transformations)
test = CamVid('./','test',transformations, transformations)
train_dataloader = data.DataLoader(train, batch_size= 4, shuffle = True,
    →num_workers=0)
val_dataloader = data.DataLoader(val, batch_size= 3, shuffle = True,
    →num_workers=0)
test_dataloader = data.DataLoader(test, batch_size= 3, shuffle = True,
    →num_workers=0)

```



```

print(len(train_dataloader), len(val_dataloader), len(test_dataloader))

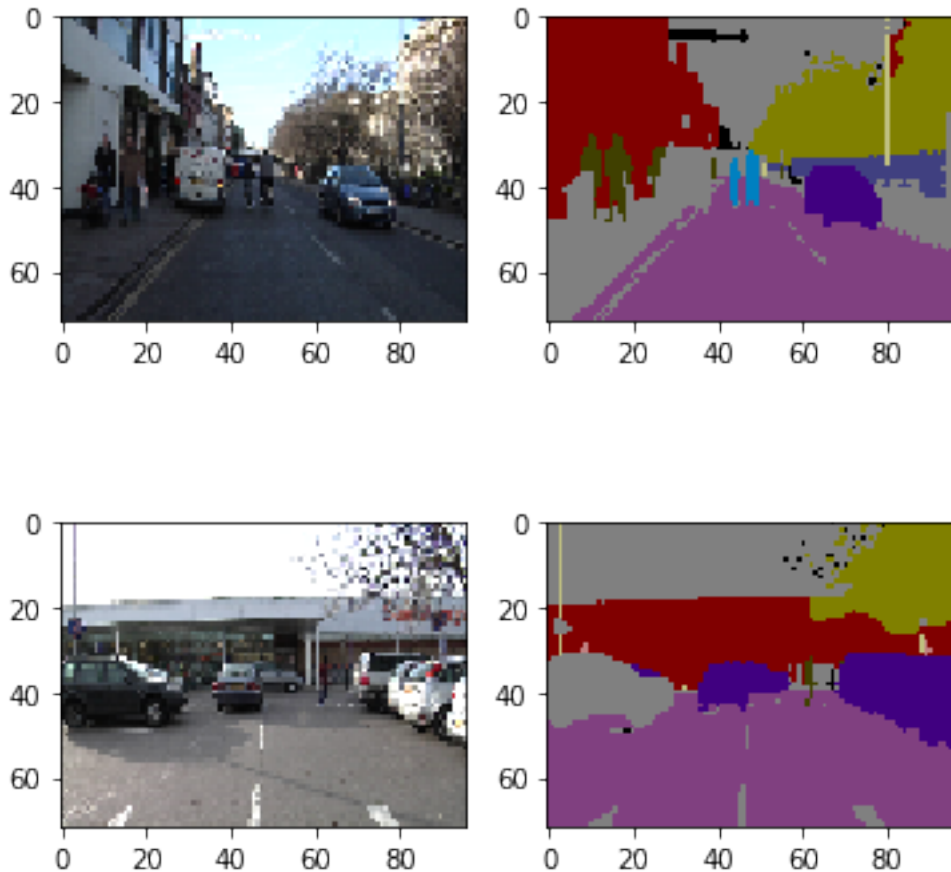
for index,[img,label] in enumerate(train_dataloader):
    print(img.size())
    print(label.size())
    mask_encoded = [label_to_rgb(label[x,:,:].numpy(), color_encoding) for x in
→range(label.shape[0])]
    for i in range(0,4):
        plt.figure()
        plt.subplot(1,2,1)
        plt.imshow(img[i].permute(1,2,0))
        plt.subplot(1,2,2)
        plt.imshow(mask_encoded[i])
    break

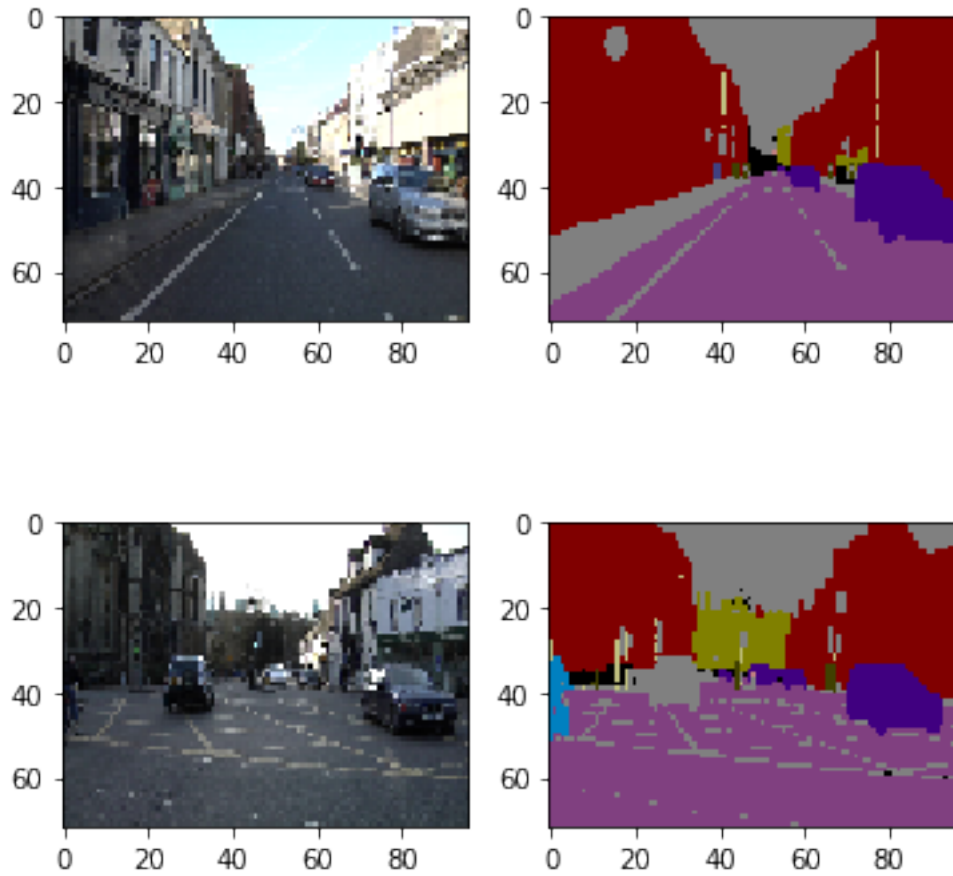
```

```

112 66 20
torch.Size([4, 3, 72, 96])
torch.Size([4, 72, 96])

```





```
[5]: # SegNet Model - CNN

DEBUG = False

vgg16_dims = [
    (64, 64, 'M'),
    (128, 128, 'M'),
    (256, 256, 256, 'M'),
    (512, 512, 512, 'M'),
    (512, 512, 512, 'M')
]

decoder_dims = [
    ('U', 512, 512, 512),
    ('U', 512, 512, 512),
    ('U', 256, 256, 256),
    ('U', 128, 128),
```

```

        ('U', 64, 64)
    ]

class SegNet(nn.Module):
    def __init__(self, input_channels, output_channels):
        super(SegNet, self).__init__()

        self.input_channels = input_channels
        self.output_channels = output_channels

        self.vgg16 = models.vgg16(pretrained=True)

        # Encoder layers

        self.encoder_conv_00 = nn.Sequential(*[
            nn.Conv2d(in_channels=self.input_channels,
                      out_channels=64,
                      kernel_size=3,
                      padding=1),
            nn.BatchNorm2d(64)
        ])
        self.encoder_conv_01 = nn.Sequential(*[
            nn.Conv2d(in_channels=64,
                      out_channels=64,
                      kernel_size=3,
                      padding=1),
            nn.BatchNorm2d(64)
        ])
        self.encoder_conv_10 = nn.Sequential(*[
            nn.Conv2d(in_channels=64,
                      out_channels=128,
                      kernel_size=3,
                      padding=1),
            nn.BatchNorm2d(128)
        ])
        self.encoder_conv_11 = nn.Sequential(*[
            nn.Conv2d(in_channels=128,
                      out_channels=128,
                      kernel_size=3,
                      padding=1),
            nn.BatchNorm2d(128)
        ])
        self.encoder_conv_20 = nn.Sequential(*[
            nn.Conv2d(in_channels=128,
                      out_channels=256,

```

```

        kernel_size=3,
        padding=1),
        nn.BatchNorm2d(256)
    ])

    self.encoder_conv_21 = nn.Sequential(*[
        nn.Conv2d(in_channels=256,
                    out_channels=256,
                    kernel_size=3,
                    padding=1),
        nn.BatchNorm2d(256)
    ])

    self.encoder_conv_22 = nn.Sequential(*[
        nn.Conv2d(in_channels=256,
                    out_channels=256,
                    kernel_size=3,
                    padding=1),
        nn.BatchNorm2d(256)
    ])

    self.encoder_conv_30 = nn.Sequential(*[
        nn.Conv2d(in_channels=256,
                    out_channels=512,
                    kernel_size=3,
                    padding=1),
        nn.BatchNorm2d(512)
    ])

    self.encoder_conv_31 = nn.Sequential(*[
        nn.Conv2d(in_channels=512,
                    out_channels=512,
                    kernel_size=3,
                    padding=1),
        nn.BatchNorm2d(512)
    ])

    self.encoder_conv_32 = nn.Sequential(*[
        nn.Conv2d(in_channels=512,
                    out_channels=512,
                    kernel_size=3,
                    padding=1),
        nn.BatchNorm2d(512)
    ])

    self.encoder_conv_40 = nn.Sequential(*[
        nn.Conv2d(in_channels=512,
                    out_channels=512,
                    kernel_size=3,
                    padding=1),
        nn.BatchNorm2d(512)
    ])

    self.encoder_conv_41 = nn.Sequential(*[

```

```

        nn.Conv2d(in_channels=512,
                  out_channels=512,
                  kernel_size=3,
                  padding=1),
        nn.BatchNorm2d(512)
    ])

    self.encoder_conv_42 = nn.Sequential(*[
        nn.Conv2d(in_channels=512,
                  out_channels=512,
                  kernel_size=3,
                  padding=1),
        nn.BatchNorm2d(512)
    ])

    self.init_vgg_weights()

    # Decoder layers

    self.decoder_convtr_42 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=512,
                          out_channels=512,
                          kernel_size=3,
                          padding=1),
        nn.BatchNorm2d(512)
    ])

    self.decoder_convtr_41 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=512,
                          out_channels=512,
                          kernel_size=3,
                          padding=1),
        nn.BatchNorm2d(512)
    ])

    self.decoder_convtr_40 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=512,
                          out_channels=512,
                          kernel_size=3,
                          padding=1),
        nn.BatchNorm2d(512)
    ])

    self.decoder_convtr_32 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=512,
                          out_channels=512,
                          kernel_size=3,
                          padding=1),
        nn.BatchNorm2d(512)
    ])

    self.decoder_convtr_31 = nn.Sequential(*[

```

```

        nn.ConvTranspose2d(in_channels=512,
                           out_channels=512,
                           kernel_size=3,
                           padding=1),
        nn.BatchNorm2d(512)
    ])

    self.decoder_convtr_30 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=512,
                           out_channels=256,
                           kernel_size=3,
                           padding=1),
        nn.BatchNorm2d(256)
    ])

    self.decoder_convtr_22 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=256,
                           out_channels=256,
                           kernel_size=3,
                           padding=1),
        nn.BatchNorm2d(256)
    ])

    self.decoder_convtr_21 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=256,
                           out_channels=256,
                           kernel_size=3,
                           padding=1),
        nn.BatchNorm2d(256)
    ])

    self.decoder_convtr_20 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=256,
                           out_channels=128,
                           kernel_size=3,
                           padding=1),
        nn.BatchNorm2d(128)
    ])

    self.decoder_convtr_11 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=128,
                           out_channels=128,
                           kernel_size=3,
                           padding=1),
        nn.BatchNorm2d(128)
    ])

    self.decoder_convtr_10 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=128,
                           out_channels=64,
                           kernel_size=3,
                           padding=1),
        nn.BatchNorm2d(64)
    ])

```

```

    ])
    self.decoder_convtr_01 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=64,
                           out_channels=64,
                           kernel_size=3,
                           padding=1),
        nn.BatchNorm2d(64)
    ])
    self.decoder_convtr_00 = nn.Sequential(*[
        nn.ConvTranspose2d(in_channels=64,
                           out_channels=self.output_channels,
                           kernel_size=3,
                           padding=1)
    ])

def forward(self, input_img):
    """
    Forward pass `input_img` through the network
    """

    # Encoder

    # Encoder Stage - 1
    dim_0 = input_img.size()
    x_00 = F.relu(self.encoder_conv_00(input_img))
    x_01 = F.relu(self.encoder_conv_01(x_00))
    x_0, indices_0 = F.max_pool2d(x_01, kernel_size=2, stride=2,
    ↪return_indices=True)

    # Encoder Stage - 2
    dim_1 = x_0.size()
    x_10 = F.relu(self.encoder_conv_10(x_0))
    x_11 = F.relu(self.encoder_conv_11(x_10))
    x_1, indices_1 = F.max_pool2d(x_11, kernel_size=2, stride=2,
    ↪return_indices=True)

    # Encoder Stage - 3
    dim_2 = x_1.size()
    x_20 = F.relu(self.encoder_conv_20(x_1))
    x_21 = F.relu(self.encoder_conv_21(x_20))
    x_22 = F.relu(self.encoder_conv_22(x_21))
    x_2, indices_2 = F.max_pool2d(x_22, kernel_size=2, stride=2,
    ↪return_indices=True)

    # Encoder Stage - 4
    dim_3 = x_2.size()

```



```

x_30 = F.relu(self.encoder_conv_30(x_2))
x_31 = F.relu(self.encoder_conv_31(x_30))
x_32 = F.relu(self.encoder_conv_32(x_31))
x_3, indices_3 = F.max_pool2d(x_32, kernel_size=2, stride=2,
↪return_indices=True)

    # Encoder Stage - 5
    dim_4 = x_3.size()
    x_40 = F.relu(self.encoder_conv_40(x_3))
    x_41 = F.relu(self.encoder_conv_41(x_40))
    x_42 = F.relu(self.encoder_conv_42(x_41))
    x_4, indices_4 = F.max_pool2d(x_42, kernel_size=2, stride=2,
↪return_indices=True)

    # Decoder

    dim_d = x_4.size()

    # Decoder Stage - 5
    x_4d = F.max_unpool2d(x_4, indices_4, kernel_size=2, stride=2,
↪output_size=dim_4)
    x_42d = F.relu(self.decoder_convtr_42(x_4d))
    x_41d = F.relu(self.decoder_convtr_41(x_42d))
    x_40d = F.relu(self.decoder_convtr_40(x_41d))
    dim_4d = x_40d.size()

    # Decoder Stage - 4
    x_3d = F.max_unpool2d(x_40d, indices_3, kernel_size=2, stride=2,
↪output_size=dim_3)
    x_32d = F.relu(self.decoder_convtr_32(x_3d))
    x_31d = F.relu(self.decoder_convtr_31(x_32d))
    x_30d = F.relu(self.decoder_convtr_30(x_31d))
    dim_3d = x_30d.size()

    # Decoder Stage - 3
    x_2d = F.max_unpool2d(x_30d, indices_2, kernel_size=2, stride=2,
↪output_size=dim_2)
    x_22d = F.relu(self.decoder_convtr_22(x_2d))
    x_21d = F.relu(self.decoder_convtr_21(x_22d))
    x_20d = F.relu(self.decoder_convtr_20(x_21d))
    dim_2d = x_20d.size()

    # Decoder Stage - 2
    x_1d = F.max_unpool2d(x_20d, indices_1, kernel_size=2, stride=2,
↪output_size=dim_1)
    x_11d = F.relu(self.decoder_convtr_11(x_1d))

```

```

x_10d = F.relu(self.decoder_convtr_10(x_11d))
dim_1d = x_10d.size()

# Decoder Stage - 1
x_0d = F.max_unpool2d(x_10d, indices_0, kernel_size=2, stride=2,
→output_size=dim_0)
x_01d = F.relu(self.decoder_convtr_01(x_0d))
x_00d = self.decoder_convtr_00(x_01d)
dim_0d = x_00d.size()

x_softmax = F.softmax(x_00d, dim=1)

if DEBUG:
    print("dim_0: {}".format(dim_0))
    print("dim_1: {}".format(dim_1))
    print("dim_2: {}".format(dim_2))
    print("dim_3: {}".format(dim_3))
    print("dim_4: {}".format(dim_4))

    print("dim_d: {}".format(dim_d))
    print("dim_4d: {}".format(dim_4d))
    print("dim_3d: {}".format(dim_3d))
    print("dim_2d: {}".format(dim_2d))
    print("dim_1d: {}".format(dim_1d))
    print("dim_0d: {}".format(dim_0d))

return x_00d, x_softmax

def init_vgg_weights(self):
    assert self.encoder_conv_00[0].weight.size() == self.vgg16.features[0].
→weight.size()
    self.encoder_conv_00[0].weight.data = self.vgg16.features[0].weight.data
    assert self.encoder_conv_00[0].bias.size() == self.vgg16.features[0].
→bias.size()
    self.encoder_conv_00[0].bias.data = self.vgg16.features[0].bias.data

    assert self.encoder_conv_01[0].weight.size() == self.vgg16.features[2].
→weight.size()
    self.encoder_conv_01[0].weight.data = self.vgg16.features[2].weight.data
    assert self.encoder_conv_01[0].bias.size() == self.vgg16.features[2].
→bias.size()
    self.encoder_conv_01[0].bias.data = self.vgg16.features[2].bias.data

```

```

        assert self.encoder_conv_10[0].weight.size() == self.vgg16.features[5].
→weight.size()
        self.encoder_conv_10[0].weight.data = self.vgg16.features[5].weight.data
        assert self.encoder_conv_10[0].bias.size() == self.vgg16.features[5].
→bias.size()
        self.encoder_conv_10[0].bias.data = self.vgg16.features[5].bias.data

        assert self.encoder_conv_11[0].weight.size() == self.vgg16.features[7].
→weight.size()
        self.encoder_conv_11[0].weight.data = self.vgg16.features[7].weight.data
        assert self.encoder_conv_11[0].bias.size() == self.vgg16.features[7].
→bias.size()
        self.encoder_conv_11[0].bias.data = self.vgg16.features[7].bias.data

        assert self.encoder_conv_20[0].weight.size() == self.vgg16.features[10].
→weight.size()
        self.encoder_conv_20[0].weight.data = self.vgg16.features[10].weight.data
        assert self.encoder_conv_20[0].bias.size() == self.vgg16.features[10].
→bias.size()
        self.encoder_conv_20[0].bias.data = self.vgg16.features[10].bias.data

        assert self.encoder_conv_21[0].weight.size() == self.vgg16.features[12].
→weight.size()
        self.encoder_conv_21[0].weight.data = self.vgg16.features[12].weight.data
        assert self.encoder_conv_21[0].bias.size() == self.vgg16.features[12].
→bias.size()
        self.encoder_conv_21[0].bias.data = self.vgg16.features[12].bias.data

        assert self.encoder_conv_22[0].weight.size() == self.vgg16.features[14].
→weight.size()
        self.encoder_conv_22[0].weight.data = self.vgg16.features[14].weight.data
        assert self.encoder_conv_22[0].bias.size() == self.vgg16.features[14].
→bias.size()
        self.encoder_conv_22[0].bias.data = self.vgg16.features[14].bias.data

        assert self.encoder_conv_30[0].weight.size() == self.vgg16.features[17].
→weight.size()
        self.encoder_conv_30[0].weight.data = self.vgg16.features[17].weight.data
        assert self.encoder_conv_30[0].bias.size() == self.vgg16.features[17].
→bias.size()
        self.encoder_conv_30[0].bias.data = self.vgg16.features[17].bias.data

        assert self.encoder_conv_31[0].weight.size() == self.vgg16.features[19].
→weight.size()
        self.encoder_conv_31[0].weight.data = self.vgg16.features[19].weight.data

```

```

        assert self.encoder_conv_31[0].bias.size() == self.vgg16.features[19].
→bias.size()
        self.encoder_conv_31[0].bias.data = self.vgg16.features[19].bias.data

        assert self.encoder_conv_32[0].weight.size() == self.vgg16.features[21].
→weight.size()
        self.encoder_conv_32[0].weight.data = self.vgg16.features[21].weight.data
        assert self.encoder_conv_32[0].bias.size() == self.vgg16.features[21].
→bias.size()
        self.encoder_conv_32[0].bias.data = self.vgg16.features[21].bias.data

        assert self.encoder_conv_40[0].weight.size() == self.vgg16.features[24].
→weight.size()
        self.encoder_conv_40[0].weight.data = self.vgg16.features[24].weight.data
        assert self.encoder_conv_40[0].bias.size() == self.vgg16.features[24].
→bias.size()
        self.encoder_conv_40[0].bias.data = self.vgg16.features[24].bias.data

        assert self.encoder_conv_41[0].weight.size() == self.vgg16.features[26].
→weight.size()
        self.encoder_conv_41[0].weight.data = self.vgg16.features[26].weight.data
        assert self.encoder_conv_41[0].bias.size() == self.vgg16.features[26].
→bias.size()
        self.encoder_conv_41[0].bias.data = self.vgg16.features[26].bias.data

        assert self.encoder_conv_42[0].weight.size() == self.vgg16.features[28].
→weight.size()
        self.encoder_conv_42[0].weight.data = self.vgg16.features[28].weight.data
        assert self.encoder_conv_42[0].bias.size() == self.vgg16.features[28].
→bias.size()
        self.encoder_conv_42[0].bias.data = self.vgg16.features[28].bias.data

```

Test for a single image

```

[6]: # Extract a single image and mask
for index,[img,label] in enumerate(train_dataloader):
    x = img
    y = label
    break

```

```

[87]: """
Train a SegNet model
Usage:
python train.py --data_root /home/SharedData/intern_sayan/PascalVOC2012/data/
→VOCdevkit/VOC2012/ \
                --train_path ImageSets/Segmentation/train.txt \

```

```

        --img_dir JPEGImages \
        --mask_dir SegmentationClass \
        --save_dir /home/SharedData/intern_sayan/PascalVOC2012/ \
        --checkpoint /home/SharedData/intern_sayan/PascalVOC2012/
→model_best.pth \
        --gpu 1
"""

from __future__ import print_function
import os
import time
import torch
from torch.utils.data import DataLoader
from torch.autograd import Variable

# Parameters
NUM_INPUT_CHANNELS = 3
NUM_OUTPUT_CHANNELS = 13

NUM_EPOCHS = 200

LEARNING_RATE = 1e-3
MOMENTUM = 0.9
BATCH_SIZE = 4

def train():
    is_better = True
    prev_loss = float('inf')
    # for continue learning
    # model.load_state_dict(torch.load("./model_best.pth"))
    model.train()

    loss_list = []
    acc_list = []
    for epoch in range(NUM_EPOCHS):
        loss_f = 0
        t_start = time.time()

        input_tensor = Variable(x)
        target_tensor = Variable(y)
        if CUDA:
            input_tensor = input_tensor.cuda()

```

```

        target_tensor = target_tensor.cuda()

    predicted_tensor, softmaxed_tensor = model(input_tensor)

    optimizer.zero_grad()
    loss = criterion(predicted_tensor, target_tensor)
    loss.backward()
    optimizer.step()

    loss_f += loss.float()

    _, prediction = torch.max(predicted_tensor, 1)
    y_true = torch.flatten(target_tensor)
    y_pred = torch.flatten(prediction)
    intersection = torch.sum(y_true * y_pred)
    acc = (2. * intersection) / (torch.sum(y_true*y_true) + torch.
→sum(y_pred*y_pred))
    acc = acc.detach().cpu().numpy()
    acc_list.append(acc)

    delta = time.time() - t_start
    is_better = loss_f < prev_loss
    loss_list.append(loss_f/len(x))
    if is_better:
        prev_loss = loss_f
        torch.save(model.state_dict(), os.path.join('.', 'model_best.pth'))

    print("Epoch #{}\tLoss: {:.8f}\t Time: {:.2f}s".format(epoch+1, loss_f,
→delta))
    return loss_list, acc_list

if __name__ == "__main__":

    CUDA = torch.cuda.is_available()

    print(CUDA)
    # train_dataloader = data.DataLoader(train, batch_size=4, shuffle = True,
→num_workers=4)

    # weights when using median frequency balancing used in SegNet paper
    # https://arxiv.org/pdf/1511.00561.pdf
    # The numbers were generated by:
    # https://github.com/yandex/segnet-torch/blob/master/datasets/camvid-gen.lua
    CAMVID_CLASS_WEIGHTS = [0.58872014284134,

```

```

0.51052379608154,
2.6966278553009,
0.45021694898605,
1.1785038709641,
0.77028578519821,
2.4782588481903,
2.5273461341858,
1.0122526884079,
3.2375309467316,
4.1312313079834,
0.3,
0]

if CUDA:
    model = SegNet(input_channels=NUM_INPUT_CHANNELS,
→output_channels=NUM_OUTPUT_CHANNELS).cuda()
    class_weights = torch.FloatTensor(CAMVID_CLASS_WEIGHTS).cuda()
    criterion = torch.nn.CrossEntropyLoss(weight=class_weights).cuda()
else:
    model = SegNet(input_channels=NUM_INPUT_CHANNELS,
→output_channels=NUM_OUTPUT_CHANNELS)
    class_weights = torch.FloatTensor(CAMVID_CLASS_WEIGHTS)
    criterion = torch.nn.CrossEntropyLoss(weight=class_weights)

optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)

loss_list = []
acc_list = []
loss_list, acc_list = train()

```

```

True
Epoch #1      Loss: 2.55754876      Time: 0.096742s
Epoch #2      Loss: 2.32248020      Time: 0.084774s
Epoch #3      Loss: 1.91855991      Time: 0.079787s
Epoch #4      Loss: 1.53522611      Time: 0.081776s
Epoch #5      Loss: 1.23839629      Time: 0.083776s
Epoch #6      Loss: 1.08583093      Time: 0.081781s
Epoch #7      Loss: 0.95990711      Time: 0.078834s
Epoch #8      Loss: 0.84045833      Time: 0.078817s
Epoch #9      Loss: 0.74290419      Time: 0.082778s
Epoch #10     Loss: 0.65871602      Time: 0.080783s
Epoch #11     Loss: 0.60670245      Time: 0.080816s
Epoch #12     Loss: 0.55186510      Time: 0.083745s
Epoch #13     Loss: 0.47650722      Time: 0.079786s
Epoch #14     Loss: 0.45593768      Time: 0.078789s
Epoch #15     Loss: 0.39415997      Time: 0.077837s
Epoch #16     Loss: 0.37400636      Time: 0.081813s

```



Epoch #17	Loss: 0.33640644	Time: 0.079787s
Epoch #18	Loss: 0.31342551	Time: 0.081816s
Epoch #19	Loss: 0.29602915	Time: 0.086768s
Epoch #20	Loss: 0.27636454	Time: 0.083809s
Epoch #21	Loss: 0.25970283	Time: 0.081781s
Epoch #22	Loss: 0.24345094	Time: 0.082803s
Epoch #23	Loss: 0.22949116	Time: 0.081814s
Epoch #24	Loss: 0.21790645	Time: 0.078788s
Epoch #25	Loss: 0.20751855	Time: 0.080784s
Epoch #26	Loss: 0.19523451	Time: 0.081780s
Epoch #27	Loss: 0.18633641	Time: 0.082814s
Epoch #28	Loss: 0.17760716	Time: 0.082778s
Epoch #29	Loss: 0.16844785	Time: 0.082778s
Epoch #30	Loss: 0.16031061	Time: 0.087765s
Epoch #31	Loss: 0.15232986	Time: 0.086764s
Epoch #32	Loss: 0.14553648	Time: 0.080784s
Epoch #33	Loss: 0.13827479	Time: 0.081780s
Epoch #34	Loss: 0.13154685	Time: 0.083775s
Epoch #35	Loss: 0.12555325	Time: 0.078820s
Epoch #36	Loss: 0.11904900	Time: 0.081816s
Epoch #37	Loss: 0.11341165	Time: 0.081808s
Epoch #38	Loss: 0.10758978	Time: 0.080781s
Epoch #39	Loss: 0.10226027	Time: 0.085739s
Epoch #40	Loss: 0.09695361	Time: 0.088763s
Epoch #41	Loss: 0.09219056	Time: 0.081781s
Epoch #42	Loss: 0.08743983	Time: 0.080785s
Epoch #43	Loss: 0.08294282	Time: 0.080822s
Epoch #44	Loss: 0.07861019	Time: 0.080781s
Epoch #45	Loss: 0.07444242	Time: 0.083793s
Epoch #46	Loss: 0.07048448	Time: 0.080784s
Epoch #47	Loss: 0.06666543	Time: 0.080784s
Epoch #48	Loss: 0.06289131	Time: 0.081781s
Epoch #49	Loss: 0.05930397	Time: 0.082810s
Epoch #50	Loss: 0.05600355	Time: 0.081795s
Epoch #51	Loss: 0.05259738	Time: 0.079798s
Epoch #52	Loss: 0.04952695	Time: 0.083776s
Epoch #53	Loss: 0.04655929	Time: 0.083774s
Epoch #54	Loss: 0.04383587	Time: 0.081262s
Epoch #55	Loss: 0.04127233	Time: 0.081781s
Epoch #56	Loss: 0.03889567	Time: 0.085771s
Epoch #57	Loss: 0.03671874	Time: 0.090757s
Epoch #58	Loss: 0.03468009	Time: 0.084805s
Epoch #59	Loss: 0.03283238	Time: 0.080783s
Epoch #60	Loss: 0.03112663	Time: 0.082817s
Epoch #61	Loss: 0.02948867	Time: 0.084773s
Epoch #62	Loss: 0.02804064	Time: 0.092453s
Epoch #63	Loss: 0.02670176	Time: 0.083776s
Epoch #64	Loss: 0.02541379	Time: 0.084774s


Epoch #65	Loss: 0.02425924	Time: 0.084779s
Epoch #66	Loss: 0.02319516	Time: 0.083797s
Epoch #67	Loss: 0.02210841	Time: 0.079792s
Epoch #68	Loss: 0.02111185	Time: 0.080783s
Epoch #69	Loss: 0.02018378	Time: 0.082779s
Epoch #70	Loss: 0.01928605	Time: 0.088762s
Epoch #71	Loss: 0.01847053	Time: 0.086767s
Epoch #72	Loss: 0.01770151	Time: 0.084773s
Epoch #73	Loss: 0.01694929	Time: 0.079792s
Epoch #74	Loss: 0.01624670	Time: 0.083809s
Epoch #75	Loss: 0.01556245	Time: 0.078789s
Epoch #76	Loss: 0.01491553	Time: 0.086876s
Epoch #77	Loss: 0.01432420	Time: 0.086768s
Epoch #78	Loss: 0.01374433	Time: 0.079786s
Epoch #79	Loss: 0.01321895	Time: 0.087764s
Epoch #80	Loss: 0.01269828	Time: 0.083774s
Epoch #81	Loss: 0.01223124	Time: 0.095731s
Epoch #82	Loss: 0.01178118	Time: 0.081781s
Epoch #83	Loss: 0.01137202	Time: 0.081781s
Epoch #84	Loss: 0.01094687	Time: 0.090758s
Epoch #85	Loss: 0.01059698	Time: 0.080828s
Epoch #86	Loss: 0.01023815	Time: 0.082778s
Epoch #87	Loss: 0.00987351	Time: 0.087767s
Epoch #88	Loss: 0.00953647	Time: 0.078790s
Epoch #89	Loss: 0.00922411	Time: 0.086737s
Epoch #90	Loss: 0.00892886	Time: 0.080750s
Epoch #91	Loss: 0.00864273	Time: 0.080784s
Epoch #92	Loss: 0.00837226	Time: 0.081746s
Epoch #93	Loss: 0.00811769	Time: 0.079819s
Epoch #94	Loss: 0.00787541	Time: 0.083777s
Epoch #95	Loss: 0.00763886	Time: 0.080818s
Epoch #96	Loss: 0.00742651	Time: 0.084771s
Epoch #97	Loss: 0.00720533	Time: 0.085771s
Epoch #98	Loss: 0.00700867	Time: 0.083775s
Epoch #99	Loss: 0.00681655	Time: 0.092753s
Epoch #100	Loss: 0.00663306	Time: 0.083776s
Epoch #101	Loss: 0.00645724	Time: 0.084773s
Epoch #102	Loss: 0.00629514	Time: 0.084062s
Epoch #103	Loss: 0.00613045	Time: 0.082851s
Epoch #104	Loss: 0.00597936	Time: 0.081829s
Epoch #105	Loss: 0.00583058	Time: 0.081806s
Epoch #106	Loss: 0.00570537	Time: 0.083773s
Epoch #107	Loss: 0.00556377	Time: 0.080783s
Epoch #108	Loss: 0.00543111	Time: 0.079842s
Epoch #109	Loss: 0.00530550	Time: 0.085768s
Epoch #110	Loss: 0.00518038	Time: 0.078821s
Epoch #111	Loss: 0.00506751	Time: 0.084774s
Epoch #112	Loss: 0.00496021	Time: 0.082778s

Epoch #113	Loss: 0.00485127	Time: 0.082747s
Epoch #114	Loss: 0.00475028	Time: 0.084772s
Epoch #115	Loss: 0.00465720	Time: 0.081786s
Epoch #116	Loss: 0.00456407	Time: 0.080783s
Epoch #117	Loss: 0.00447794	Time: 0.083775s
Epoch #118	Loss: 0.00438806	Time: 0.082793s
Epoch #119	Loss: 0.00429752	Time: 0.078799s
Epoch #120	Loss: 0.00420985	Time: 0.082778s
Epoch #121	Loss: 0.00412936	Time: 0.086768s
Epoch #122	Loss: 0.00405786	Time: 0.079787s
Epoch #123	Loss: 0.00398485	Time: 0.080796s
Epoch #124	Loss: 0.00391647	Time: 0.082778s
Epoch #125	Loss: 0.00385709	Time: 0.081770s
Epoch #126	Loss: 0.00379302	Time: 0.082790s
Epoch #127	Loss: 0.00372220	Time: 0.080815s
Epoch #128	Loss: 0.00365963	Time: 0.081781s
Epoch #129	Loss: 0.00359425	Time: 0.085771s
Epoch #130	Loss: 0.00353090	Time: 0.083776s
Epoch #131	Loss: 0.00347060	Time: 0.080783s
Epoch #132	Loss: 0.00341160	Time: 0.079787s
Epoch #133	Loss: 0.00335626	Time: 0.079786s
Epoch #134	Loss: 0.00330193	Time: 0.084772s
Epoch #135	Loss: 0.00325376	Time: 0.082827s
Epoch #136	Loss: 0.00319987	Time: 0.079785s
Epoch #137	Loss: 0.00314772	Time: 0.083744s
Epoch #138	Loss: 0.00309894	Time: 0.080817s
Epoch #139	Loss: 0.00306179	Time: 0.083807s
Epoch #140	Loss: 0.00301197	Time: 0.078821s
Epoch #141	Loss: 0.00296639	Time: 0.079785s
Epoch #142	Loss: 0.00292266	Time: 0.083803s
Epoch #143	Loss: 0.00288073	Time: 0.080814s
Epoch #144	Loss: 0.00284072	Time: 0.079820s
Epoch #145	Loss: 0.00280011	Time: 0.082780s
Epoch #146	Loss: 0.00276002	Time: 0.082755s
Epoch #147	Loss: 0.00272262	Time: 0.078789s
Epoch #148	Loss: 0.00268664	Time: 0.078833s
Epoch #149	Loss: 0.00264995	Time: 0.077792s
Epoch #150	Loss: 0.00261384	Time: 0.080812s
Epoch #151	Loss: 0.00257895	Time: 0.078821s
Epoch #152	Loss: 0.00254511	Time: 0.083787s
Epoch #153	Loss: 0.00251233	Time: 0.078822s
Epoch #154	Loss: 0.00248032	Time: 0.088762s
Epoch #155	Loss: 0.00244892	Time: 0.079787s
Epoch #156	Loss: 0.00241803	Time: 0.081781s
Epoch #157	Loss: 0.00238798	Time: 0.078832s
Epoch #158	Loss: 0.00235874	Time: 0.080783s
Epoch #159	Loss: 0.00232989	Time: 0.083807s
Epoch #160	Loss: 0.00230183	Time: 0.082744s

Epoch #161	Loss: 0.00227423	Time: 0.087734s
Epoch #162	Loss: 0.00224787	Time: 0.081815s
Epoch #163	Loss: 0.00222185	Time: 0.078758s
Epoch #164	Loss: 0.00219617	Time: 0.081782s
Epoch #165	Loss: 0.00217079	Time: 0.078789s
Epoch #166	Loss: 0.00214699	Time: 0.078789s
Epoch #167	Loss: 0.00212278	Time: 0.083774s
Epoch #168	Loss: 0.00209864	Time: 0.080780s
Epoch #169	Loss: 0.00207498	Time: 0.081781s
Epoch #170	Loss: 0.00205225	Time: 0.078822s
Epoch #171	Loss: 0.00202994	Time: 0.080783s
Epoch #172	Loss: 0.00200800	Time: 0.080783s
Epoch #173	Loss: 0.00198646	Time: 0.082779s
Epoch #174	Loss: 0.00196560	Time: 0.082778s
Epoch #175	Loss: 0.00194497	Time: 0.087766s
Epoch #176	Loss: 0.00192446	Time: 0.081813s
Epoch #177	Loss: 0.00190439	Time: 0.079787s
Epoch #178	Loss: 0.00188466	Time: 0.080804s
Epoch #179	Loss: 0.00186683	Time: 0.081781s
Epoch #180	Loss: 0.00184656	Time: 0.086768s
Epoch #181	Loss: 0.00182857	Time: 0.079755s
Epoch #182	Loss: 0.00181027	Time: 0.079786s
Epoch #183	Loss: 0.00179571	Time: 0.085804s
Epoch #184	Loss: 0.00177862	Time: 0.081804s
Epoch #185	Loss: 0.00176072	Time: 0.080816s
Epoch #186	Loss: 0.00174336	Time: 0.080817s
Epoch #187	Loss: 0.00172666	Time: 0.082813s
Epoch #188	Loss: 0.00170992	Time: 0.077755s
Epoch #189	Loss: 0.00169344	Time: 0.080758s
Epoch #190	Loss: 0.00167738	Time: 0.077791s
Epoch #191	Loss: 0.00166144	Time: 0.079785s
Epoch #192	Loss: 0.00164748	Time: 0.083792s
Epoch #193	Loss: 0.00163225	Time: 0.080816s
Epoch #194	Loss: 0.00161725	Time: 0.081782s
Epoch #195	Loss: 0.00160575	Time: 0.086760s
Epoch #196	Loss: 0.00159144	Time: 0.083777s
Epoch #197	Loss: 0.00157649	Time: 0.080752s
Epoch #198	Loss: 0.00156168	Time: 0.087765s
Epoch #199	Loss: 0.00154735	Time: 0.082826s
Epoch #200	Loss: 0.00153315	Time: 0.077825s

```
[97]: # Learning curve
x1 = range(0, NUM_EPOCHS)
y1 = acc_list
plt.figure
plt.plot(x1, y1, 'b-', label="train_acc")
plt.xlabel('Test accuracy/loss vs. epoches')
```

```
plt.ylabel('Test accuracy/loss')
x2 = range(0, NUM_EPOCHS)
y2 = loss_list
plt.plot(x2, y2, 'y-', label="train_loss")
plt.legend(loc="center right")
plt.show()
```



```
[27]: # test

SAVED_MODEL_PATH = "./model_best.pth"
OUTPUT_DIR = "./result_image"

def validate():
    model.eval()

    input_tensor = torch.autograd.Variable(x)
    target_tensor = torch.autograd.Variable(y)

    if CUDA:
        input_tensor = input_tensor.cuda()
        target_tensor = target_tensor.cuda()

    predicted_tensor, softmaxed_tensor = model(input_tensor)
    loss = criterion(predicted_tensor, target_tensor)
```

```

for idx, predicted_mask in enumerate(softmaxed_tensor):

    target_mask = target_tensor[idx]
    input_image = input_tensor[idx]
    fig = plt.figure()

    a = fig.add_subplot(1,3,1)
    plt.imshow(input_image.cpu().permute(1,2,0))
    a.set_title('Input Image')

    a = fig.add_subplot(1,3,2)
    predicted_mx = predicted_mask.detach().cpu().numpy()
    predicted_mx = predicted_mx.argmax(axis=0)
    predicted_rgb = label_to_rgb(predicted_mx, color_encoding)
    plt.imshow(predicted_rgb)
    a.set_title('Predicted Mask')

    a = fig.add_subplot(1,3,3)
    target_mx = target_mask.detach().cpu().numpy()
    target_rgb = label_to_rgb(target_mx, color_encoding)
    plt.imshow(target_rgb)
    a.set_title('Ground Truth')

if __name__ == "__main__":

    CUDA = torch.cuda.is_available()

    if CUDA:
        model = SegNet(input_channels=NUM_INPUT_CHANNELS,
                        output_channels=NUM_OUTPUT_CHANNELS).cuda()





        class_weights = torch.FloatTensor(CAMVID_CLASS_WEIGHTS).cuda()
        criterion = torch.nn.CrossEntropyLoss(weight=class_weights).cuda()
    else:
        model = SegNet(input_channels=NUM_INPUT_CHANNELS,
                        output_channels=NUM_OUTPUT_CHANNELS)

        class_weights = torch.FloatTensor(CAMVID_CLASS_WEIGHTS).cuda()
        criterion = torch.nn.CrossEntropyLoss(weight=class_weights)

    model.load_state_dict(torch.load(SAVED_MODEL_PATH))

    validate()

```





Test for whole dataset

```
[8]: """
Train a SegNet model
Usage:
python train.py --data_root /home/SharedData/intern_sayan/PascalVOC2012/data/
→VOCdevkit/VOC2012/ \
    --train_path ImageSets/Segmentation/train.txt \
    --img_dir JPEGImages \
    --mask_dir SegmentationClass \
    --save_dir /home/SharedData/intern_sayan/PascalVOC2012/ \
    --checkpoint /home/SharedData/intern_sayan/PascalVOC2012/
→model_best.pth \
    --gpu 1
"""

from __future__ import print_function
import os
import time
import torch
from torch.utils.data import DataLoader
from torch.autograd import Variable

# Parameters
NUM_INPUT_CHANNELS = 3
NUM_OUTPUT_CHANNELS = 13

NUM_EPOCHS = 160

LEARNING_RATE = 1e-3
MOMENTUM = 0.9
BATCH_SIZE = 4

def train():
    is_better = True
    prev_loss = float('inf')
    # model.load_state_dict(torch.load("./model_best.pth"))
    model.train()

    loss_list = []
    acc_list = []
    for epoch in range(NUM_EPOCHS):
        loss_f = 0
```

```

t_start = time.time()

for index,[img,label] in enumerate(train_dataloader):
    input_tensor = Variable(img)
    target_tensor = Variable(label)
    if CUDA:
        input_tensor = input_tensor.cuda()
        target_tensor = target_tensor.cuda()

    predicted_tensor, softmaxed_tensor = model(input_tensor)

    optimizer.zero_grad()
    loss = criterion(predicted_tensor, target_tensor)
    loss.backward()
    optimizer.step()

    loss_f += loss.float()

    delta = time.time() - t_start
    is_better = loss_f < prev_loss

    if is_better:
        prev_loss = loss_f
    if index%40== 0:
        print("Epoch #{}\tTotal loss: {:.8f}\t Time: {:.2f}s".
→format(epoch+1, loss_f , delta))
    if index == (len(train_dataloader) - 1):
        loss_list.append(loss_f)
        _, prediction = torch.max(predicted_tensor,1)
        y_true = torch.flatten(target_tensor)
        y_prep = torch.flatten(prediction)
        intersection = torch.sum(y_true * y_prep)
        acc = (2. * intersection) / (torch.sum(y_true*y_true) + torch.
→sum(y_prep*y_prep))
        acc = acc.detach().cpu().numpy()
        acc_list.append(acc)
        torch.save(model.state_dict(), os.path.join('/content/drive/My Drive/
→segnet', "model_best.pth"))
    return loss_list,acc_list

if __name__ == "__main__":

    CUDA = torch.cuda.is_available()

```

```

# weights when using median frequency balancing used in SegNet paper
# https://arxiv.org/pdf/1511.00561.pdf
# The numbers were generated by:
# https://github.com/yandex/segnet-torch/blob/master/datasets/camvid-gen.lua
CAMVID_CLASS_WEIGHTS = [0.58872014284134,
                        0.51052379608154,
                        2.6966278553009,
                        0.45021694898605,
                        1.1785038709641,
                        0.77028578519821,
                        2.4782588481903,
                        2.5273461341858,
                        1.0122526884079,
                        3.2375309467316,
                        4.1312313079834,
                        0.3,
                        0]

if CUDA:
    model = SegNet(input_channels=NUM_INPUT_CHANNELS,
→output_channels=NUM_OUTPUT_CHANNELS).cuda()
    class_weights = torch.FloatTensor(CAMVID_CLASS_WEIGHTS).cuda()
    criterion = torch.nn.CrossEntropyLoss(weight=class_weights).cuda()
else:
    model = SegNet(input_channels=NUM_INPUT_CHANNELS,
→output_channels=NUM_OUTPUT_CHANNELS)
    class_weights = torch.FloatTensor(CAMVID_CLASS_WEIGHTS)
    criterion = torch.nn.CrossEntropyLoss(weight=class_weights)

optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)

loss_list = []
acc_list = []
loss_list, acc_list= train()

```

Epoch #1	Total loss: 2.79073524	Time: 3.158261s
Epoch #1	Total loss: 52.70711517	Time: 120.173725s
Epoch #1	Total loss: 88.79529572	Time: 225.196719s
Epoch #2	Total loss: 0.89505726	Time: 0.196335s
Epoch #2	Total loss: 34.81172180	Time: 8.891248s
Epoch #2	Total loss: 67.23095703	Time: 17.123669s
Epoch #3	Total loss: 0.66244602	Time: 0.206405s
Epoch #3	Total loss: 30.05274200	Time: 7.692185s
Epoch #3	Total loss: 57.62046814	Time: 15.973802s
Epoch #4	Total loss: 0.74270427	Time: 0.294773s

Epoch #4	Total loss: 28.73457336	Time: 7.664445s
Epoch #4	Total loss: 55.50716400	Time: 15.843583s
Epoch #5	Total loss: 0.51564145	Time: 0.194166s
Epoch #5	Total loss: 25.89620972	Time: 7.567474s
Epoch #5	Total loss: 51.41121674	Time: 15.462481s
Epoch #6	Total loss: 0.64950240	Time: 0.193268s
Epoch #6	Total loss: 27.27328300	Time: 7.767581s
Epoch #6	Total loss: 54.92087936	Time: 15.842180s
Epoch #7	Total loss: 0.48998925	Time: 0.186098s
Epoch #7	Total loss: 26.42895889	Time: 7.957027s
Epoch #7	Total loss: 52.58362579	Time: 16.471228s
Epoch #8	Total loss: 0.74209774	Time: 0.223622s
Epoch #8	Total loss: 27.88312340	Time: 7.908955s
Epoch #8	Total loss: 53.17484665	Time: 16.348693s
Epoch #9	Total loss: 0.53831977	Time: 0.201347s
Epoch #9	Total loss: 24.60177231	Time: 7.974913s
Epoch #9	Total loss: 49.43677139	Time: 16.215892s
Epoch #10	Total loss: 0.72697365	Time: 0.412176s
Epoch #10	Total loss: 28.23324776	Time: 7.981323s
Epoch #10	Total loss: 53.09757996	Time: 16.176852s
Epoch #11	Total loss: 0.68504590	Time: 0.200847s
Epoch #11	Total loss: 24.53111076	Time: 7.780811s
Epoch #11	Total loss: 49.10469437	Time: 16.369132s
Epoch #12	Total loss: 0.61855805	Time: 0.407448s
Epoch #12	Total loss: 22.39270401	Time: 8.153446s
Epoch #12	Total loss: 44.23672104	Time: 16.618912s
Epoch #13	Total loss: 0.42640942	Time: 0.374184s
Epoch #13	Total loss: 22.01772499	Time: 8.056992s
Epoch #13	Total loss: 44.43247986	Time: 16.463574s
Epoch #14	Total loss: 0.46345830	Time: 0.465798s
Epoch #14	Total loss: 20.59221840	Time: 8.092975s
Epoch #14	Total loss: 39.98138809	Time: 16.474354s
Epoch #15	Total loss: 0.59995335	Time: 0.205059s
Epoch #15	Total loss: 19.70584297	Time: 7.886991s
Epoch #15	Total loss: 38.87735748	Time: 16.212556s
Epoch #16	Total loss: 0.49734887	Time: 0.245718s
Epoch #16	Total loss: 21.05565643	Time: 7.759240s
Epoch #16	Total loss: 40.55717850	Time: 15.930613s
Epoch #17	Total loss: 0.42771584	Time: 0.251478s
Epoch #17	Total loss: 18.42227745	Time: 7.792841s
Epoch #17	Total loss: 37.36229324	Time: 15.811461s
Epoch #18	Total loss: 0.45413458	Time: 0.204554s
Epoch #18	Total loss: 17.72665977	Time: 7.821945s
Epoch #18	Total loss: 36.39314270	Time: 16.289325s
Epoch #19	Total loss: 0.33608621	Time: 0.200009s
Epoch #19	Total loss: 17.17278099	Time: 7.769358s
Epoch #19	Total loss: 34.61000061	Time: 15.870885s
Epoch #20	Total loss: 0.54773474	Time: 0.809238s

Epoch #20	Total loss: 16.87650681	Time: 8.347552s
Epoch #20	Total loss: 34.31000137	Time: 16.543852s
Epoch #21	Total loss: 0.39975718	Time: 0.192823s
Epoch #21	Total loss: 17.36980247	Time: 7.503309s
Epoch #21	Total loss: 33.26158524	Time: 15.483444s
Epoch #22	Total loss: 0.43039480	Time: 0.183169s
Epoch #22	Total loss: 17.80486107	Time: 7.606499s
Epoch #22	Total loss: 33.52763748	Time: 15.504840s
Epoch #23	Total loss: 0.38034090	Time: 0.206482s
Epoch #23	Total loss: 15.01726723	Time: 7.570601s
Epoch #23	Total loss: 30.50909233	Time: 15.484264s
Epoch #24	Total loss: 0.35522732	Time: 0.176549s
Epoch #24	Total loss: 15.95325756	Time: 7.603338s
Epoch #24	Total loss: 30.76111984	Time: 15.536719s
Epoch #25	Total loss: 0.50881809	Time: 0.186021s
Epoch #25	Total loss: 13.93685150	Time: 7.468513s
Epoch #25	Total loss: 28.21644211	Time: 15.367339s
Epoch #26	Total loss: 0.33947930	Time: 0.236910s
Epoch #26	Total loss: 14.77280617	Time: 7.518933s
Epoch #26	Total loss: 30.83484268	Time: 15.434861s
Epoch #27	Total loss: 0.30337632	Time: 0.180359s
Epoch #27	Total loss: 15.20313644	Time: 7.381647s
Epoch #27	Total loss: 30.13105011	Time: 15.254979s
Epoch #28	Total loss: 0.30573216	Time: 0.383163s
Epoch #28	Total loss: 14.43111229	Time: 7.611681s
Epoch #28	Total loss: 27.69159317	Time: 15.530839s
Epoch #29	Total loss: 0.42956463	Time: 0.178907s
Epoch #29	Total loss: 16.00214386	Time: 7.504362s
Epoch #29	Total loss: 35.28704834	Time: 15.438761s
Epoch #30	Total loss: 0.40119320	Time: 0.196252s
Epoch #30	Total loss: 16.87842751	Time: 7.429649s
Epoch #30	Total loss: 34.23440552	Time: 15.355723s
Epoch #31	Total loss: 0.36637092	Time: 0.175045s
Epoch #31	Total loss: 15.41750717	Time: 7.447023s
Epoch #31	Total loss: 29.84320831	Time: 15.270514s
Epoch #32	Total loss: 0.37262589	Time: 0.251554s
Epoch #32	Total loss: 18.47575951	Time: 7.463608s
Epoch #32	Total loss: 35.51514435	Time: 15.385399s
Epoch #33	Total loss: 0.48846874	Time: 0.181816s
Epoch #33	Total loss: 15.25479889	Time: 7.427532s
Epoch #33	Total loss: 29.82523537	Time: 15.336196s
Epoch #34	Total loss: 0.34673735	Time: 0.567045s
Epoch #34	Total loss: 15.30326080	Time: 7.707755s
Epoch #34	Total loss: 29.16212463	Time: 15.552392s
Epoch #35	Total loss: 0.35141057	Time: 0.176940s
Epoch #35	Total loss: 12.55681705	Time: 7.372004s
Epoch #35	Total loss: 24.52288818	Time: 15.237844s
Epoch #36	Total loss: 0.35156521	Time: 0.186207s

Epoch #36	Total loss: 12.37071228	Time: 7.451581s
Epoch #36	Total loss: 24.92478943	Time: 15.299366s
Epoch #37	Total loss: 0.24568884	Time: 0.159665s
Epoch #37	Total loss: 11.10037899	Time: 7.310243s
Epoch #37	Total loss: 21.78655052	Time: 15.403590s
Epoch #38	Total loss: 0.32508337	Time: 0.410890s
Epoch #38	Total loss: 10.96766853	Time: 7.559850s
Epoch #38	Total loss: 22.72075653	Time: 15.405246s
Epoch #39	Total loss: 0.47703439	Time: 0.214341s
Epoch #39	Total loss: 15.94780636	Time: 7.354910s
Epoch #39	Total loss: 29.26250839	Time: 15.193760s
Epoch #40	Total loss: 0.22967039	Time: 0.315267s
Epoch #40	Total loss: 12.00907803	Time: 7.464911s
Epoch #40	Total loss: 23.06777954	Time: 15.131840s
Epoch #41	Total loss: 0.25352502	Time: 0.290790s
Epoch #41	Total loss: 10.80625916	Time: 7.428892s
Epoch #41	Total loss: 21.83430672	Time: 15.281844s
Epoch #42	Total loss: 0.21151462	Time: 0.209105s
Epoch #42	Total loss: 9.63568592	Time: 7.386911s
Epoch #42	Total loss: 20.20479012	Time: 15.234480s
Epoch #43	Total loss: 0.20672701	Time: 0.194951s
Epoch #43	Total loss: 10.16579914	Time: 7.435892s
Epoch #43	Total loss: 20.61584091	Time: 15.224571s
Epoch #44	Total loss: 0.38105777	Time: 0.207663s
Epoch #44	Total loss: 13.04780197	Time: 7.379904s
Epoch #44	Total loss: 24.87861252	Time: 15.302292s
Epoch #45	Total loss: 0.30065984	Time: 0.357611s
Epoch #45	Total loss: 11.59985924	Time: 7.557507s
Epoch #45	Total loss: 23.56425095	Time: 15.311113s
Epoch #46	Total loss: 0.16784479	Time: 0.302363s
Epoch #46	Total loss: 11.22827148	Time: 7.508335s
Epoch #46	Total loss: 22.39805984	Time: 15.342592s
Epoch #47	Total loss: 0.24182041	Time: 0.225774s
Epoch #47	Total loss: 9.81881714	Time: 7.416037s
Epoch #47	Total loss: 18.77026749	Time: 15.102608s
Epoch #48	Total loss: 0.17978711	Time: 0.256214s
Epoch #48	Total loss: 8.79106426	Time: 7.461378s
Epoch #48	Total loss: 17.33232117	Time: 15.292384s
Epoch #49	Total loss: 0.22315247	Time: 0.274798s
Epoch #49	Total loss: 8.21559811	Time: 7.547099s
Epoch #49	Total loss: 16.21633720	Time: 15.336759s
Epoch #50	Total loss: 0.27824315	Time: 0.217484s
Epoch #50	Total loss: 8.04436970	Time: 7.410898s
Epoch #50	Total loss: 15.86022568	Time: 15.185818s
Epoch #51	Total loss: 0.16689165	Time: 0.399042s
Epoch #51	Total loss: 7.41523933	Time: 7.616218s
Epoch #51	Total loss: 14.46589088	Time: 15.448824s
Epoch #52	Total loss: 0.13586032	Time: 0.246193s

Epoch #52	Total loss: 7.54118919	Time: 7.394165s
Epoch #52	Total loss: 15.09787941	Time: 15.238204s
Epoch #53	Total loss: 0.25433078	Time: 0.307779s
Epoch #53	Total loss: 7.55999756	Time: 7.478351s
Epoch #53	Total loss: 15.45461273	Time: 15.222769s
Epoch #54	Total loss: 0.20231220	Time: 0.205498s
Epoch #54	Total loss: 7.75809097	Time: 7.371165s
Epoch #54	Total loss: 15.36217403	Time: 15.131016s
Epoch #55	Total loss: 0.18836553	Time: 0.163409s
Epoch #55	Total loss: 8.41700840	Time: 7.297954s
Epoch #55	Total loss: 16.26173019	Time: 15.169752s
Epoch #56	Total loss: 0.16360240	Time: 0.177052s
Epoch #56	Total loss: 6.98355722	Time: 7.357804s
Epoch #56	Total loss: 13.88643646	Time: 15.124029s
Epoch #57	Total loss: 0.18606143	Time: 0.182432s
Epoch #57	Total loss: 6.43459654	Time: 7.376367s
Epoch #57	Total loss: 13.31478500	Time: 15.269836s
Epoch #58	Total loss: 0.11146921	Time: 0.186228s
Epoch #58	Total loss: 6.88897562	Time: 7.425875s
Epoch #58	Total loss: 13.26956081	Time: 15.198689s
Epoch #59	Total loss: 0.13871455	Time: 0.184455s
Epoch #59	Total loss: 6.21291351	Time: 7.405334s
Epoch #59	Total loss: 12.23571110	Time: 15.188972s
Epoch #60	Total loss: 0.14386167	Time: 0.450551s
Epoch #60	Total loss: 5.64130449	Time: 7.629508s
Epoch #60	Total loss: 11.18725586	Time: 15.504489s
Epoch #61	Total loss: 0.10836481	Time: 0.173029s
Epoch #61	Total loss: 7.52743864	Time: 7.313670s
Epoch #61	Total loss: 17.30227661	Time: 15.135767s
Epoch #62	Total loss: 0.26762527	Time: 0.181603s
Epoch #62	Total loss: 22.52822495	Time: 7.273490s
Epoch #62	Total loss: 43.45865250	Time: 14.953768s
Epoch #63	Total loss: 0.51760143	Time: 0.171182s
Epoch #63	Total loss: 17.58595657	Time: 7.349326s
Epoch #63	Total loss: 33.94009018	Time: 15.092829s
Epoch #64	Total loss: 0.35991395	Time: 0.178358s
Epoch #64	Total loss: 11.49263859	Time: 7.540679s
Epoch #64	Total loss: 23.05109978	Time: 15.431317s
Epoch #65	Total loss: 0.21813080	Time: 0.172369s
Epoch #65	Total loss: 9.47453499	Time: 7.271712s
Epoch #65	Total loss: 18.58236313	Time: 14.933399s
Epoch #66	Total loss: 0.16488995	Time: 0.176511s
Epoch #66	Total loss: 8.44415474	Time: 7.289031s
Epoch #66	Total loss: 16.52661133	Time: 14.936495s
Epoch #67	Total loss: 0.32048601	Time: 0.164366s
Epoch #67	Total loss: 7.87690687	Time: 7.600685s
Epoch #67	Total loss: 15.52514362	Time: 15.897325s
Epoch #68	Total loss: 0.17968494	Time: 0.201982s



Epoch #68	Total loss: 7.06397963	Time: 7.829511s
Epoch #68	Total loss: 14.39227962	Time: 16.232228s
Epoch #69	Total loss: 0.12908733	Time: 0.200681s
Epoch #69	Total loss: 6.91664743	Time: 7.809687s
Epoch #69	Total loss: 13.49976921	Time: 16.162857s
Epoch #70	Total loss: 0.14612368	Time: 0.186625s
Epoch #70	Total loss: 6.34257507	Time: 7.838933s
Epoch #70	Total loss: 12.30103397	Time: 16.361441s
Epoch #71	Total loss: 0.12584981	Time: 0.199649s
Epoch #71	Total loss: 5.97572517	Time: 7.818001s
Epoch #71	Total loss: 11.50465107	Time: 16.120608s
Epoch #72	Total loss: 0.13791458	Time: 0.186275s
Epoch #72	Total loss: 6.45688009	Time: 7.766879s
Epoch #72	Total loss: 13.68082523	Time: 15.913530s
Epoch #73	Total loss: 0.12559366	Time: 0.184407s
Epoch #73	Total loss: 5.91043234	Time: 7.788450s
Epoch #73	Total loss: 11.82211971	Time: 16.118456s
Epoch #74	Total loss: 0.16804776	Time: 0.193503s
Epoch #74	Total loss: 6.15995646	Time: 7.837401s
Epoch #74	Total loss: 11.69609928	Time: 16.120982s
Epoch #75	Total loss: 0.11361320	Time: 0.178457s
Epoch #75	Total loss: 5.55689192	Time: 7.877071s
Epoch #75	Total loss: 10.67031193	Time: 16.120674s
Epoch #76	Total loss: 0.14774768	Time: 0.191190s
Epoch #76	Total loss: 5.29548931	Time: 7.972161s
Epoch #76	Total loss: 10.43679619	Time: 16.220072s
Epoch #77	Total loss: 0.11461330	Time: 0.206772s
Epoch #77	Total loss: 5.49796677	Time: 8.113926s
Epoch #77	Total loss: 10.92245960	Time: 16.437419s
Epoch #78	Total loss: 0.14723153	Time: 0.187198s
Epoch #78	Total loss: 5.60879469	Time: 7.996922s
Epoch #78	Total loss: 11.59295177	Time: 16.173334s
Epoch #79	Total loss: 0.12484452	Time: 0.190699s
Epoch #79	Total loss: 5.53596020	Time: 7.811692s
Epoch #79	Total loss: 11.24941635	Time: 16.435960s
Epoch #80	Total loss: 0.13350672	Time: 0.206242s
Epoch #80	Total loss: 5.09477234	Time: 8.274406s
Epoch #80	Total loss: 10.53414726	Time: 16.941365s
Epoch #81	Total loss: 0.13318625	Time: 0.998603s
Epoch #81	Total loss: 4.80633163	Time: 9.221313s
Epoch #81	Total loss: 9.81284904	Time: 18.052896s
Epoch #82	Total loss: 0.13543141	Time: 0.214278s
Epoch #82	Total loss: 5.03738070	Time: 8.206817s
Epoch #82	Total loss: 10.10075855	Time: 17.022867s
Epoch #83	Total loss: 0.12292805	Time: 0.190571s
Epoch #83	Total loss: 5.05745935	Time: 8.407451s
Epoch #83	Total loss: 10.05570412	Time: 17.432245s
Epoch #84	Total loss: 0.13884325	Time: 0.213753s

Epoch #84	Total loss: 4.78088951	Time: 8.823522s
Epoch #84	Total loss: 9.61703491	Time: 17.890411s
Epoch #85	Total loss: 0.14235201	Time: 0.216515s
Epoch #85	Total loss: 5.50257206	Time: 8.216386s
Epoch #85	Total loss: 11.13757229	Time: 16.861474s
Epoch #86	Total loss: 0.17665485	Time: 0.198692s
Epoch #86	Total loss: 5.60017729	Time: 8.278206s
Epoch #86	Total loss: 10.95112419	Time: 16.779225s
Epoch #87	Total loss: 0.11005562	Time: 0.575612s
Epoch #87	Total loss: 5.21832132	Time: 8.257239s
Epoch #87	Total loss: 10.09813404	Time: 16.660527s
Epoch #88	Total loss: 0.12479895	Time: 0.190830s
Epoch #88	Total loss: 4.72096825	Time: 7.934304s
Epoch #88	Total loss: 9.61163616	Time: 16.254495s
Epoch #89	Total loss: 0.11709055	Time: 0.219411s
Epoch #89	Total loss: 4.77782297	Time: 8.265243s
Epoch #89	Total loss: 9.21133041	Time: 16.867628s
Epoch #90	Total loss: 0.12548842	Time: 0.197069s
Epoch #90	Total loss: 4.33230782	Time: 8.027630s
Epoch #90	Total loss: 9.22784138	Time: 16.446584s
Epoch #91	Total loss: 0.16294156	Time: 0.188218s
Epoch #91	Total loss: 5.37274742	Time: 8.122585s
Epoch #91	Total loss: 14.91723156	Time: 16.648874s
Epoch #92	Total loss: 0.17041677	Time: 0.203429s
Epoch #92	Total loss: 10.00752544	Time: 8.107420s
Epoch #92	Total loss: 19.35751343	Time: 16.313532s
Epoch #93	Total loss: 0.18508707	Time: 0.200658s
Epoch #93	Total loss: 7.72218990	Time: 8.009469s
Epoch #93	Total loss: 14.94191647	Time: 16.457192s
Epoch #94	Total loss: 0.16343530	Time: 0.215199s
Epoch #94	Total loss: 6.72229052	Time: 7.919339s
Epoch #94	Total loss: 12.45981407	Time: 16.503967s
Epoch #95	Total loss: 0.11938729	Time: 0.179026s
Epoch #95	Total loss: 5.25648975	Time: 7.856662s
Epoch #95	Total loss: 10.56264210	Time: 16.176284s
Epoch #96	Total loss: 0.09528979	Time: 0.201320s
Epoch #96	Total loss: 4.82281065	Time: 7.894253s
Epoch #96	Total loss: 9.55266285	Time: 16.367084s
Epoch #97	Total loss: 0.09730206	Time: 0.205249s
Epoch #97	Total loss: 4.46385193	Time: 8.077306s
Epoch #97	Total loss: 8.76303101	Time: 16.444807s
Epoch #98	Total loss: 0.11357366	Time: 0.206909s
Epoch #98	Total loss: 4.25098944	Time: 7.919779s
Epoch #98	Total loss: 8.37553883	Time: 16.390149s
Epoch #99	Total loss: 0.08575506	Time: 0.200410s
Epoch #99	Total loss: 4.39531326	Time: 7.763487s
Epoch #99	Total loss: 8.36770916	Time: 16.033583s
Epoch #100	Total loss: 0.11028317	Time: 0.195043s

Epoch #100	Total loss: 4.01414824	Time: 7.716946s
Epoch #100	Total loss: 7.89959192	Time: 15.818573s
Epoch #101	Total loss: 0.06925537	Time: 0.209399s
Epoch #101	Total loss: 4.05975199	Time: 8.210423s
Epoch #101	Total loss: 8.19833374	Time: 16.761515s
Epoch #102	Total loss: 0.11370429	Time: 0.180517s
Epoch #102	Total loss: 3.84588909	Time: 7.970098s
Epoch #102	Total loss: 7.72584820	Time: 16.293585s
Epoch #103	Total loss: 0.09912010	Time: 0.219903s
Epoch #103	Total loss: 3.71868277	Time: 7.673444s
Epoch #103	Total loss: 7.46599388	Time: 15.969617s
Epoch #104	Total loss: 0.09330463	Time: 0.191726s
Epoch #104	Total loss: 4.83966923	Time: 7.628542s
Epoch #104	Total loss: 9.68351936	Time: 15.529906s
Epoch #105	Total loss: 0.10503995	Time: 0.190225s
Epoch #105	Total loss: 4.38699436	Time: 7.598211s
Epoch #105	Total loss: 8.68557167	Time: 15.764383s
Epoch #106	Total loss: 0.10274321	Time: 0.192111s
Epoch #106	Total loss: 4.19416666	Time: 7.730830s
Epoch #106	Total loss: 8.40557194	Time: 15.737056s
Epoch #107	Total loss: 0.12908451	Time: 0.204476s
Epoch #107	Total loss: 4.17168570	Time: 7.658205s
Epoch #107	Total loss: 8.10866928	Time: 15.749087s
Epoch #108	Total loss: 0.09755995	Time: 0.181942s
Epoch #108	Total loss: 3.94235444	Time: 7.564597s
Epoch #108	Total loss: 8.00729656	Time: 15.582460s
Epoch #109	Total loss: 0.11044141	Time: 0.190091s
Epoch #109	Total loss: 13.40043068	Time: 7.675172s
Epoch #109	Total loss: 31.28546715	Time: 15.494831s
Epoch #110	Total loss: 0.32479265	Time: 0.182283s
Epoch #110	Total loss: 12.57219315	Time: 7.551666s
Epoch #110	Total loss: 21.88905716	Time: 15.514877s
Epoch #111	Total loss: 0.22977939	Time: 0.175295s
Epoch #111	Total loss: 7.57708549	Time: 7.549559s
Epoch #111	Total loss: 15.15770435	Time: 15.599338s
Epoch #112	Total loss: 0.17266135	Time: 0.202516s
Epoch #112	Total loss: 6.02539110	Time: 7.554028s
Epoch #112	Total loss: 11.30610657	Time: 15.721478s
Epoch #113	Total loss: 0.13099469	Time: 0.188635s
Epoch #113	Total loss: 4.63828754	Time: 7.729051s
Epoch #113	Total loss: 9.46577930	Time: 15.921593s
Epoch #114	Total loss: 0.10014349	Time: 0.195898s
Epoch #114	Total loss: 4.33158636	Time: 7.892416s
Epoch #114	Total loss: 8.74739075	Time: 15.976326s
Epoch #115	Total loss: 0.09349649	Time: 0.182276s
Epoch #115	Total loss: 4.27468443	Time: 7.671766s
Epoch #115	Total loss: 8.18968582	Time: 15.806005s
Epoch #116	Total loss: 0.07677813	Time: 0.181106s

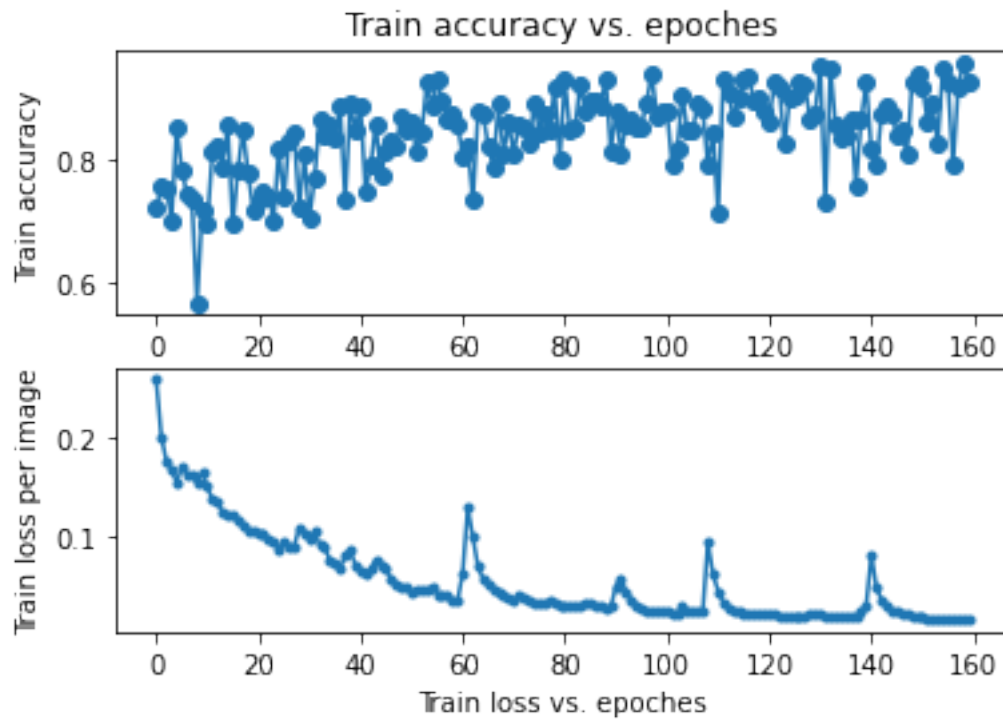
Epoch #116	Total loss: 4.05060005	Time: 7.427478s
Epoch #116	Total loss: 7.99387598	Time: 15.386001s
Epoch #117	Total loss: 0.07024020	Time: 0.189835s
Epoch #117	Total loss: 3.93690133	Time: 7.465838s
Epoch #117	Total loss: 7.82861185	Time: 15.244781s
Epoch #118	Total loss: 0.09780766	Time: 0.186970s
Epoch #118	Total loss: 3.76240015	Time: 7.478864s
Epoch #118	Total loss: 7.38265753	Time: 15.531018s
Epoch #119	Total loss: 0.08668015	Time: 0.192884s
Epoch #119	Total loss: 3.64459348	Time: 7.914633s
Epoch #119	Total loss: 7.21307611	Time: 15.758930s
Epoch #120	Total loss: 0.08306948	Time: 0.193478s
Epoch #120	Total loss: 3.66238999	Time: 7.587528s
Epoch #120	Total loss: 7.50117874	Time: 15.780506s
Epoch #121	Total loss: 0.07867153	Time: 0.184141s
Epoch #121	Total loss: 3.62266803	Time: 7.415211s
Epoch #121	Total loss: 7.35357761	Time: 15.166422s
Epoch #122	Total loss: 0.08435074	Time: 0.184600s
Epoch #122	Total loss: 3.52651024	Time: 7.412456s
Epoch #122	Total loss: 7.11246109	Time: 15.218717s
Epoch #123	Total loss: 0.06576300	Time: 0.185416s
Epoch #123	Total loss: 3.45994210	Time: 7.361278s
Epoch #123	Total loss: 6.95947599	Time: 15.138936s
Epoch #124	Total loss: 0.06928811	Time: 0.178518s
Epoch #124	Total loss: 3.28147244	Time: 7.470210s
Epoch #124	Total loss: 6.58287954	Time: 15.293926s
Epoch #125	Total loss: 0.09120315	Time: 0.169430s
Epoch #125	Total loss: 3.29278970	Time: 7.349483s
Epoch #125	Total loss: 6.51277781	Time: 14.893424s
Epoch #126	Total loss: 0.05617649	Time: 0.180432s
Epoch #126	Total loss: 3.29943252	Time: 7.376539s
Epoch #126	Total loss: 6.45654058	Time: 15.133730s
Epoch #127	Total loss: 0.06876162	Time: 0.180581s
Epoch #127	Total loss: 3.05317283	Time: 7.429797s
Epoch #127	Total loss: 6.13008213	Time: 15.427504s
Epoch #128	Total loss: 0.08044403	Time: 0.171240s
Epoch #128	Total loss: 3.49249458	Time: 7.275644s
Epoch #128	Total loss: 6.70160627	Time: 15.028857s
Epoch #129	Total loss: 0.10716099	Time: 0.185010s
Epoch #129	Total loss: 3.78790641	Time: 7.410002s
Epoch #129	Total loss: 7.60303354	Time: 15.212306s
Epoch #130	Total loss: 0.06745390	Time: 0.175373s
Epoch #130	Total loss: 3.69922352	Time: 7.325103s
Epoch #130	Total loss: 7.28573370	Time: 15.042457s
Epoch #131	Total loss: 0.08976488	Time: 0.181720s
Epoch #131	Total loss: 3.65254831	Time: 7.398531s
Epoch #131	Total loss: 7.50702429	Time: 15.243118s
Epoch #132	Total loss: 0.10312302	Time: 0.177512s

Epoch #132	Total loss: 3.44609499	Time: 7.300535s
Epoch #132	Total loss: 6.78670454	Time: 14.981159s
Epoch #133	Total loss: 0.08445737	Time: 0.170148s
Epoch #133	Total loss: 3.23099113	Time: 7.400976s
Epoch #133	Total loss: 6.52785683	Time: 15.889786s
Epoch #134	Total loss: 0.08857644	Time: 0.200073s
Epoch #134	Total loss: 3.26301575	Time: 8.091557s
Epoch #134	Total loss: 6.41276503	Time: 16.651599s
Epoch #135	Total loss: 0.06457938	Time: 0.184986s
Epoch #135	Total loss: 3.01787257	Time: 7.895486s
Epoch #135	Total loss: 6.23055172	Time: 15.520697s
Epoch #136	Total loss: 0.07028575	Time: 0.190239s
Epoch #136	Total loss: 3.05905724	Time: 8.112459s
Epoch #136	Total loss: 6.17266941	Time: 16.159835s
Epoch #137	Total loss: 0.09520821	Time: 0.207305s
Epoch #137	Total loss: 3.14060783	Time: 8.048711s
Epoch #137	Total loss: 6.32169914	Time: 15.848579s
Epoch #138	Total loss: 0.05729937	Time: 0.209660s
Epoch #138	Total loss: 3.08809304	Time: 8.531389s
Epoch #138	Total loss: 6.26960516	Time: 16.431709s
Epoch #139	Total loss: 0.06828906	Time: 0.191434s
Epoch #139	Total loss: 3.36435676	Time: 7.942687s
Epoch #139	Total loss: 7.57109404	Time: 16.241220s
Epoch #140	Total loss: 0.09568826	Time: 0.223541s
Epoch #140	Total loss: 4.56288767	Time: 8.507626s
Epoch #140	Total loss: 9.23435211	Time: 16.422458s
Epoch #141	Total loss: 0.20904543	Time: 0.224469s
Epoch #141	Total loss: 10.90330505	Time: 8.068180s
Epoch #141	Total loss: 26.55052757	Time: 16.161376s
Epoch #142	Total loss: 0.19079681	Time: 0.204668s
Epoch #142	Total loss: 8.78553867	Time: 8.109647s
Epoch #142	Total loss: 16.50287247	Time: 16.205364s
Epoch #143	Total loss: 0.15287465	Time: 0.214536s
Epoch #143	Total loss: 6.20199633	Time: 8.100671s
Epoch #143	Total loss: 12.17472649	Time: 15.937210s
Epoch #144	Total loss: 0.09948855	Time: 0.212908s
Epoch #144	Total loss: 5.01955605	Time: 8.430067s
Epoch #144	Total loss: 10.25812244	Time: 16.367984s
Epoch #145	Total loss: 0.09879492	Time: 0.230086s
Epoch #145	Total loss: 4.49010515	Time: 8.426300s
Epoch #145	Total loss: 8.64378643	Time: 17.428938s
Epoch #146	Total loss: 0.09164838	Time: 0.207455s
Epoch #146	Total loss: 4.11086273	Time: 8.292027s
Epoch #146	Total loss: 7.91448164	Time: 16.366685s
Epoch #147	Total loss: 0.10102090	Time: 0.219768s
Epoch #147	Total loss: 3.65038824	Time: 8.104664s
Epoch #147	Total loss: 7.40128326	Time: 15.909150s
Epoch #148	Total loss: 0.09974888	Time: 0.205407s

Epoch #148	Total loss: 3.67055154	Time: 7.943307s
Epoch #148	Total loss: 7.25462818	Time: 15.619517s
Epoch #149	Total loss: 0.08021094	Time: 0.206821s
Epoch #149	Total loss: 3.44124651	Time: 8.006305s
Epoch #149	Total loss: 6.88147926	Time: 15.660192s
Epoch #150	Total loss: 0.07705156	Time: 0.202717s
Epoch #150	Total loss: 3.26406097	Time: 7.864746s
Epoch #150	Total loss: 6.40975380	Time: 15.561658s
Epoch #151	Total loss: 0.11058713	Time: 0.214940s
Epoch #151	Total loss: 3.13103604	Time: 8.233959s
Epoch #151	Total loss: 6.24022961	Time: 16.837000s
Epoch #152	Total loss: 0.06229006	Time: 0.207310s
Epoch #152	Total loss: 2.95474696	Time: 8.366466s
Epoch #152	Total loss: 6.02444601	Time: 16.672201s
Epoch #153	Total loss: 0.07468464	Time: 0.218290s
Epoch #153	Total loss: 2.87955213	Time: 8.156497s
Epoch #153	Total loss: 5.88207150	Time: 15.927642s
Epoch #154	Total loss: 0.08137368	Time: 0.213184s
Epoch #154	Total loss: 3.07300568	Time: 8.329954s
Epoch #154	Total loss: 6.02743101	Time: 16.273614s
Epoch #155	Total loss: 0.11463949	Time: 0.204260s
Epoch #155	Total loss: 2.94165254	Time: 7.978554s
Epoch #155	Total loss: 5.98151016	Time: 15.910996s
Epoch #156	Total loss: 0.06781179	Time: 0.201692s
Epoch #156	Total loss: 2.92896199	Time: 8.231076s
Epoch #156	Total loss: 5.76968288	Time: 16.321428s
Epoch #157	Total loss: 0.06371322	Time: 0.210765s
Epoch #157	Total loss: 2.88053584	Time: 8.138907s
Epoch #157	Total loss: 5.71014929	Time: 17.127388s
Epoch #158	Total loss: 0.07307822	Time: 0.210353s
Epoch #158	Total loss: 2.87134433	Time: 7.986340s
Epoch #158	Total loss: 5.83848810	Time: 15.723507s
Epoch #159	Total loss: 0.08712765	Time: 0.212622s
Epoch #159	Total loss: 2.87067533	Time: 8.170836s
Epoch #159	Total loss: 5.73263741	Time: 16.266184s
Epoch #160	Total loss: 0.08165640	Time: 0.201157s
Epoch #160	Total loss: 2.77239561	Time: 8.234527s
Epoch #160	Total loss: 5.40727901	Time: 16.042111s

```
[10]: x1 = range(0, NUM_EPOCHS)
      y1 = acc_list
      plt.subplot(2, 1, 1)
      plt.plot(x1, y1, 'o-')
      plt.title('Train accuracy vs. epoches')
      plt.ylabel('Train accuracy')
      x2 = range(0, NUM_EPOCHS)
      y2 = np.array(loss_list)/(len(train_dataloader)*4)
```

```
plt.subplot(2, 1, 2)
plt.plot(x2, y2, '-.')
plt.xlabel('Train loss vs. epoches')
plt.ylabel('Train loss per image')
plt.show()
#
```



```
[11]: # test
SAVED_MODEL_PATH = "/content/drive/My Drive/segnet/model_best.pth"
OUTPUT_DIR = "/content/drive/My Drive/segnet/result_image"

# Parameters
NUM_INPUT_CHANNELS = 3
NUM_OUTPUT_CHANNELS = 13

MOMENTUM = 0.9
BATCH_SIZE = 4
CAMVID_CLASS_WEIGHTS = [0.58872014284134,
                        0.51052379608154,
                        2.6966278553009,
```

```

0.45021694898605,
1.1785038709641,
0.77028578519821,
2.4782588481903,
2.5273461341858,
1.0122526884079,
3.2375309467316,
4.1312313079834,
0.3,
0]
def validate():
    model.eval()
    acc_list_val = []
    loss_list_val = []
    for batch_idx, [img, label] in enumerate(val_dataloader):
        input_tensor = torch.autograd.Variable(img)
        target_tensor = torch.autograd.Variable(label)

        if CUDA:
            input_tensor = input_tensor.cuda()
            target_tensor = target_tensor.cuda()

        predicted_tensor, softmaxed_tensor = model(input_tensor)
        loss = criterion(predicted_tensor, target_tensor)
        loss_list_val.append(loss.float())

        _, prediction = torch.max(predicted_tensor, 1)
        y_true = torch.flatten(target_tensor)
        y_prep = torch.flatten(prediction)
        intersection = torch.sum(y_true * y_prep)
        acc = (2. * intersection) / (torch.sum(y_true*y_true) + torch.
→sum(y_prep*y_prep))
        acc = acc.detach().cpu().numpy()
        acc_list_val.append(acc)
        print("Processd #{} batch out of {}/{}".format(batch_idx+1, batch_idx+1,
→8))

    if batch_idx == 8:
        print("Visualize the result for 3 images in last batch")
        for idx, predicted_mask in enumerate(softmaxed_tensor):

            target_mask = target_tensor[idx]
            input_image = input_tensor[idx]
            fig = plt.figure()

            a = fig.add_subplot(1,3,1)
            plt.imshow(input_image.cpu().permute(1,2,0))

```



```

        a.set_title('Input Image')

        a = fig.add_subplot(1,3,2)
        predicted_mx = predicted_mask.detach().cpu().numpy()
        predicted_mx = predicted_mx.argmax(axis=0)
        predicted_rgb = label_to_rgb(predicted_mx, color_encoding)
        plt.imshow(predicted_rgb)
        a.set_title('Predicted Mask')

        a = fig.add_subplot(1,3,3)
        target_mx = target_mask.detach().cpu().numpy()
        target_rgb = label_to_rgb(target_mx, color_encoding)
        plt.imshow(target_rgb)
        a.set_title('Ground Truth')

        fig.savefig(os.path.join(OUTPUT_DIR, "prediction_{}_{}.png".
→format(batch_idx+1, idx)))
        break
    return loss_list_val, acc_list_val
if __name__ == "__main__":

    CUDA = torch.cuda.is_available()
    if CUDA:
        model = SegNet(input_channels=NUM_INPUT_CHANNELS,
                        output_channels=NUM_OUTPUT_CHANNELS).cuda()

        class_weights = torch.FloatTensor(CAMVID_CLASS_WEIGHTS).cuda()
        criterion = torch.nn.CrossEntropyLoss(weight=class_weights).cuda()
    else:
        model = SegNet(input_channels=NUM_INPUT_CHANNELS,
                        output_channels=NUM_OUTPUT_CHANNELS)

        class_weights = torch.FloatTensor(CAMVID_CLASS_WEIGHTS).cuda()
        criterion = torch.nn.CrossEntropyLoss(weight=class_weights)

    model.load_state_dict(torch.load(SAVED_MODEL_PATH))

    loss_list_val = []
    acc_list_val = []
    loss_list_val, acc_list_val = validate()




```


```

Processd #1 batch out of 1/8
Processd #2 batch out of 2/8
Processd #3 batch out of 3/8
Processd #4 batch out of 4/8
Processd #5 batch out of 5/8

```

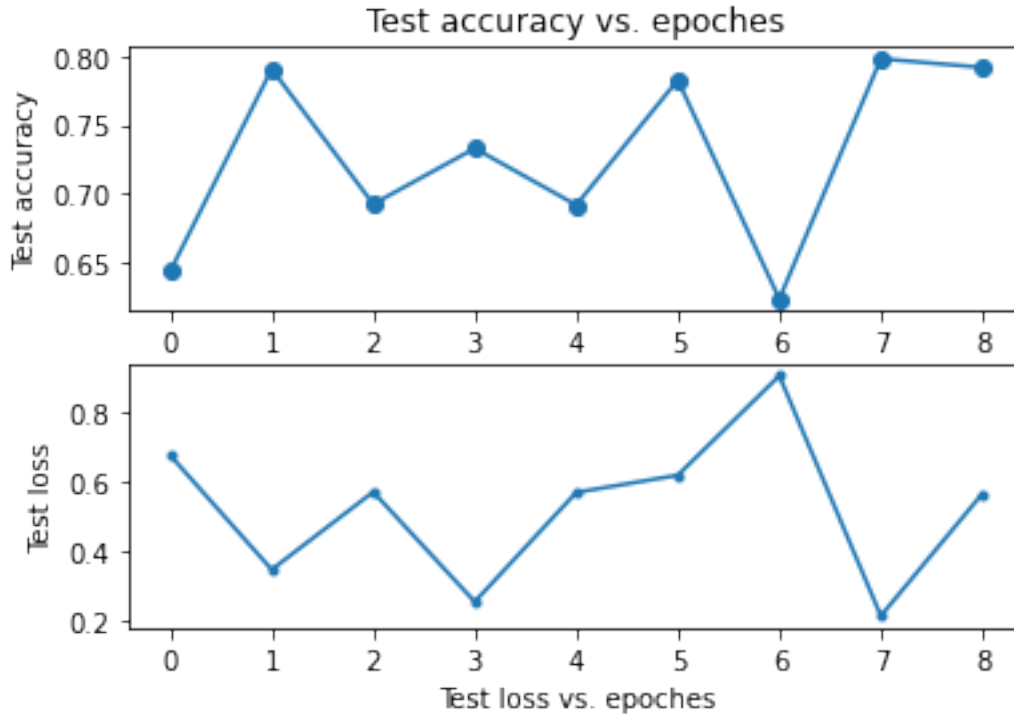
Processd #6 batch out of 6/8  
Processd #7 batch out of 7/8  
Processd #8 batch out of 8/8  
Processd #9 batch out of 9/8  
Visualize the result for 3 images in last batch





```
[15]: x1 = range(0, len(acc_list_val))
y1 = acc_list_val
plt.subplot(2, 1, 1)
plt.plot(x1, y1, 'o-')
plt.title('Test accuracy vs. epoches')
plt.ylabel('Test accuracy')
x2 = range(0, len(acc_list_val))
y2 = np.array(loss_list_val)/4

plt.subplot(2, 1, 2)
plt.plot(x2, y2, '-.')
plt.xlabel('Test loss vs. epoches')
plt.ylabel('Test loss')
plt.show()
#
mean_acc = np.sum(y1)/len(y1)
mean_loss = np.sum(y2)/len(y2)
print('Mean accuracy on validation set is {:.2f}'.format(mean_acc))
print('Mean loss on validation set is {:.2f}'.format(mean_loss))
```



Mean accuracy on validation set is 0.727296

Mean loss on validation set is 0.523363

## 0.7 Conclusion

From single image training and test result, we validate the usability of the segnet. And we can see the model overfitted on single image and result “perfect” predicted mask.

For whole dataset training, we conducted 160 epochs. The trend of the learning curve and accuracy curve are plausible (one converges to the apparent error one improves overall), but still because our selection of the learning rate and varying image distributions the curves are with some fluctuations.

For the validation part, we draw 8 mini batches from the validation set. The overall mean accuracy among all test batches turns to be 72.7%, comparing to the correspond value given in the paper is 61.9%, but that doesn’t mean our result is well performed since the mean loss is not that low comparing to the overfitting single image loss. But by comparing the visualized results, we can manually check the real performance. The predicted mask and ground truth image are matched well, even same on a lot of detailed features. So my inference that the lower mean acc may be caused by some specific objects mis matching, e.g. in comparison images 3 and 4 the vehicle on the sideway are entirely not clustered at all.

There are still lots of improvements we can do on our reproduction project. Since we don’t have official source code, we implemented lots of functions ourselves, some of them are time and space inefficient, which boosts the training time epically when the image resolution is high. That’s also the reason we can’t eventually train on the full-size data. Secondly, we didn’t use advanced ML

techniques in this project such as early stopping, varying learning rate and etc. But if those can be added we can see the pros and cons on using these methods.

For illumination invariance image processing implementing those two image processing methods, the RGB images can be successfully transformed into illumination-invariant grayscale images. It can be concluded from the processed images that the influence of varying illumination has been eliminated. For example, the intensity of a whole wall in the gray image will keep consistent even though the wall has half of it covered by the shadow. Since the original RGB images have three channels as the input of the neural network, the processed grayscale images need to be repeated for three times to act as the three channels of the input to fit the pre-trained VGG16 network. As for the code, the image processing function serves as part of the transformation function. Once a batch of images were chosen to be trained or tested, they would be transformed into illumination-invariant images. Unfortunately, the processed images could not be trained successfully because the prediction array always kept being 'nan' from the first epoch for some unknown reasons, which caused back propagation impossible.

## 0.8 Reference

1. G. D. Finlayson and S. D. Hordley, "Color constancy at a pixel," *Journal of the Optical Society of America. A, Optics, Image Science, and Vision*, vol. 18, no. 2, pp. 253–264, 2001.
2. W. Maddern, A. Stewart, C. McManus, B. Upcroft, W. Churchill, and P. Newman, "Illumination invariant imaging: Applications in robust vision-based localisation, mapping and classification for autonomous vehicles," in *Proc. Int. Conf. on Robotics and Automation*, vol. 2, 2014, p. 3.
3. J. A. Alvarez and A. Lopez, "Road detection based on illuminant invariance," *IEEE Trans. on Intelligent Transportation Systems*, vol. 12, no. 1, pp. 184–193, 2011.
4. V. Badrinarayanan, A. Kendall, and R. Cipolla, "Segnet: A deep convolutional encoder-decoder architecture for image segmentation," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 39, no. 12, pp. 2481–2495, 2017.
5. K. Simonyan and A. Zisserman, "Very deep convolutional networks for large-scale image recognition," *CoRR*, vol. abs/1409.1556, 2014.
6. A. Krizhevsky, I. Sutskever, and G. Hinton, "Imagenet classification with deep convolutional neural networks," in *Advances in Neural Information Processing Systems 25*, 2012.
7. <https://towardsdatascience.com/adam-latest-trends-in-deep-learning-optimization-6be9a291375c>